

Shortest Paths, and Dijkstra's Algorithm: Overview

- Graphs with lengths/weights/costs on edges.
- Shortest paths in edge-weighted graphs
- Dijkstra's classic algorithm for computing single-source shortest paths.

Graphs with edge “length” (or “weight/cost”)

An **edge-weighted directed graph**, $G = (V, E, w)$, has a **length/weight/cost** function, $w : E \rightarrow \mathbb{N}$, which maps each edge $(u, v) \in E$ to a non-negative integer “length” (or “weight”, or “cost”): $w(u, v) \in \mathbb{N}$.

We can **extend** the “length” function w to a function $w : V \times V \rightarrow \mathbb{N} \cup \{\infty\}$, by letting $w(u, u) = 0$, for all $u \in V$, and letting $w(u, v) = \infty$ for all $(u, v) \notin E$.

Consider a directed path:

$$x_0 e_1 x_1 e_2 \dots e_n x_n$$

from $u = x_0 \in V$ to $v = x_n \in V$, in graph $G = (V, E, w)$. The **length** of this path is defined to be: $\sum_{i=1}^n w(x_{i-1}, x_i)$.

Graphs with edge “length” (or “weight/cost”)

An **edge-weighted directed graph**, $G = (V, E, w)$, has a **length/weight/cost** function, $w : E \rightarrow \mathbb{N}$, which maps each edge $(u, v) \in E$ to a non-negative integer “length” (or “weight”, or “cost”): $w(u, v) \in \mathbb{N}$.

We can **extend** the “length” function w to a function $w : V \times V \rightarrow \mathbb{N} \cup \{\infty\}$, by letting $w(u, u) = 0$, for all $u \in V$, and letting $w(u, v) = \infty$ for all $(u, v) \notin E$.

Consider a directed path:

$$x_0 e_1 x_1 e_2 \dots e_n x_n$$

from $u = x_0 \in V$ to $v = x_n \in V$, in graph $G = (V, E, w)$. The **length** of this path is defined to be: $\sum_{i=1}^n w(x_{i-1}, x_i)$.

Question: Given G and a pair of vertices $u, v \in V$, how do we compute the length of the **shortest path** from u to v ?

Dijkstra's single-source shortest-path algorithm

Input: Edge-weighted graph, $G = (V, E, w)$, with (*extended*) weight function $w : V \times V \rightarrow \mathbb{N}$, and a source vertex $s \in V$.

Output: Function $L : V \rightarrow \mathbb{N} \cup \{\infty\}$, such that for all $v \in V$, $L(v)$ is the length of the shortest path from s to v in G .

Algorithm:

Initialize: $S := \{s\}$; $L(s) := 0$;

Initialize: $L(v) := w(s, v)$, for all $v \in V - \{s\}$;

while ($S \neq V$) **do**

$u := \arg \min_{z \in V - S} \{L(z)\}$

$S := S \cup \{u\}$

for all $v \in V - S$ such that $(u, v) \in E$ **do**

$L(v) := \min\{L(v), L(u) + w(u, v)\}$

end for

end while

Output function $L(\cdot)$.

Why does Dijkstra's algorithm work?

Claim: The While loop of Dijkstra's algorithm maintains the following **invariant** properties of the function L and the set S :

1. $\forall v \in S, L(v)$ is the shortest path length from s to v in G .
2. $\forall v \in V - S, L(v)$ is the length of the shortest path from s to v which uses only vertices in $S \cup \{v\}$.
3. For all $u \in S$ and $v \in V - S, L(u) \leq L(v)$.

Note that the three invariants hold after initialization, just prior to the first iteration of the while loop.

The claim follows once we prove (on board) that **if** the invariants hold just prior to a while loop iteration **then** they hold just after.

Since each iteration adds one vertex to S , it follows that the algorithm halts, at which point $S = V$, and thus, by invariant (1.), the function $L : V \rightarrow \mathbb{N} \cup \{\infty\}$ is the correct answer.

Remarks on Dijkstra's Algorithm

- If Dijkstra's algorithm is implemented naively, it has running time $O(n^2)$, where $n = |V|$.
- With clever data structures (e.g., so called “Fibonacci Heaps”) Dijkstra's algorithm can be implemented much more efficiently: essentially in time $O(m + n \log n)$ where, $n = |V|$ and $m = |E|$.

This increased efficiency can make a **big difference** on huge “sparse” graphs, where m is much smaller than n^2 (e.g., when out-degree is a fixed constant, $m \in O(n)$).

- Dijkstra's algorithm can be augmented to also output a description of a shortest path from the source vertex s to every other vertex v .

We will not describe these extensions, and we will certainly not assume that you know them.

Graph Colouring

Graph Colouring

Suppose we have k distinct colours with which to colour the vertices of a graph. Let $[k] = \{1, \dots, k\}$. For an undirected graph, $G = (V, E)$, an admissible vertex **k -colouring** of G is a function $c: V \rightarrow [k]$, such that for all $u, v \in V$, if $\{u, v\} \in E$ then $c(u) \neq c(v)$.

For an integer $k \geq 1$, we say an undirected graph $G = (V, E)$ is **k -colourable** if there exists a k -colouring of G .

The **chromatic number** of G , denoted $\chi(G)$, is the *smallest positive integer* k , such that G is k -colourable.

Some observations about Graph colouring

- Note that any graph G with n vertices is n -colourable.
- The **n -Clique**, K_n , i.e., the complete graph on n vertices, has chromatic number $\chi(K_n) = n$. All its vertices must get assigned different colours in any admissible colouring.
- The **clique number**, $\omega(G)$, of a graph G is the maximum positive integer $r \geq 1$, such that K_r is a subgraph of G .
- Note that for all graphs G , $\omega(G) \leq \chi(G)$: if G has an r -clique then it is not $(r - 1)$ -colorable.
- However, in general, $\omega(G) \neq \chi(G)$. For instance, The 5-cycle, C_5 , has $\omega(C_5) = 2 < \chi(C_5) = 3$.

More observations about colouring

- As already mentioned, any bipartite graph is 2-colourable. Indeed, that is an equivalent definition of being bipartite.
- More generally, a graph G is k -colourable precisely if it is k -partite, meaning its vertices can be partitioned into k disjoint sets such that all edges of the graph are between nodes in different parts.

Algorithms/complexity of colouring graphs

To determine whether a n -vertex graph $G = (V, E)$ is k -colourable by “*brute force*”, we could try all possible colourings of n nodes with k colours.

Difficulty: There are k^n such k -colouring functions $c : V \rightarrow [k]$.

Question: Is there an efficient (polynomial time) algorithm for determining whether a given graph G is k -colourable?

Algorithms/complexity of colouring graphs

To determine whether a n -vertex graph $G = (V, E)$ is k -colourable by “*brute force*”, we could try all possible colourings of n nodes with k colours.

Difficulty: There are k^n such k -colouring functions $c : V \rightarrow [k]$.

Question: Is there an efficient (polynomial time) algorithm for determining whether a given graph G is k -colourable?

Answer: No, no generally efficient (polynomial time) algorithm is known, and even the problem of determining whether a given graph is 3-colourable is **NP-complete**. (Even approximating the chromatic number of a given graph is NP-hard.)

In practice, there are heuristic algorithms that do obtain good colourings for many classes of graphs.

Applications of Graph Colouring (many)

Final Exam Scheduling

- There are n courses, $\{1, \dots, n\}$.
- Some courses have the same students registered for both, so their exams can't be scheduled at the same time.
- Let $G = (\{1, \dots, n\}, E)$ be a graph such that $\{i, j\} \in E$ if and only if $i \neq j$ and courses i and j have a student in common.
- **Question:** What is the minimum number of exam time slots needed to schedule all n exams?
- **Answer:** This is precisely the chromatic number $\chi(G)$ of G .

Furthermore, a k -colouring of G yields an *admissible schedule* of exams into k time slots, allowing all students to attend all their exams, as long as different “colors” are scheduled in disjoint time slots.