

SQL: Structured Query Language

Part III

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use **NOT IN**, **NOT EXISTS** and **NOT UNIQUE**.
- Also available: *op ANY*, *op ALL* >, <, =, ≥, ≤, ≠
- *Find sailors whose rating is greater than some sailor called Jim:*

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ANY
```

```
(SELECT S2.rating  
FROM   Sailors S2  
WHERE  S2.sname='Jim')
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

If the subquery returns an **empty set**, comparison returns **FALSE** ²

More on Set-Comparison Operators

- Find sailors whose rating is greater than every sailor called Jim.*

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ALL
```

```
(SELECT  S2.rating  
FROM    Sailors S2  
WHERE   S2.sname='Jim')
```

If the subquery returns an empty set, comparison returns TRUE

More on Set-Comparison Operators

- Find sailors with highest rating.*

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Jim	1	17
22	Fred	7	50
3	Nancy	1	21

```
SELECT *  
FROM Sailors S  
WHERE S.rating >= ALL (SELECT S2.rating  
                        FROM Sailors S2)
```

*Note: **IN** equivalent to **= ANY**
NOT IN equivalent to **<> ALL***

Division in SQL

Find sailors who've reserved all boats.

$$\rho \text{ (} Temp\text{sids, (} \pi_{sid,bid} \text{Reserves) / (} \pi_{bid} \text{Boats))}$$
$$\pi_{sname} \text{ (} Temp\text{sids} \bowtie \text{Sailors)}$$

Let's think in a complementary way.

Sailors without any unreserved boat.

Division in SQL

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S
```

Sailors S such that ...

```
WHERE NOT EXISTS (
```

```
SELECT B.bid
```

there is no boat B

```
FROM Boats B
```

without ...

```
WHERE NOT EXISTS ( SELECT R.bid
```

```
FROM Reserves R
```

a Reserves tuple showing S reserved B

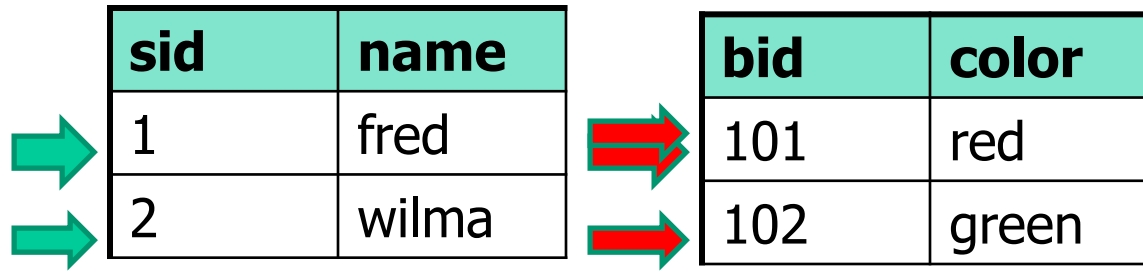
```
WHERE R.bid=B.bid
```

(for which there is no reservation tuple)

```
AND R.sid=S.sid ) )
```

Subquery in the green box:

The boats that are **not reserved** by the given sailor!



Answer:

sid
Fred

SELECT S.sname
FROM Sailors S

WHERE NOT EXISTS (SELECT B.bid
FROM Boats B
WHERE NOT EXISTS (SELECT R.bid
FROM Reserves R
WHERE R.bid=B.bid
AND R.sid=S.sid))

sid	bid
1	101
1	102
2	102

Division using EXCEPT

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S  
WHERE NOT EXISTS
```

```
((SELECT B.bid  
FROM Boats B)
```

EXCEPT

```
(SELECT R.bid  
FROM Reserves R  
WHERE R.sid=S.sid))
```

name
fred

bid
101
102

bid
102

bid
101
102

sid	name
1	fred
2	wilma

bid	color
101	red
102	green

sid	bid
1	101
1	102
2	102

Aggregate Operators

- Significant extension of relational algebra.

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A)
MIN (A)

single column

Aggregate Operators

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Fred'
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=7
```

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=7
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Mary	1	17
22	Fred	7	50
3	Nancy	2	21

Aggregate Operators

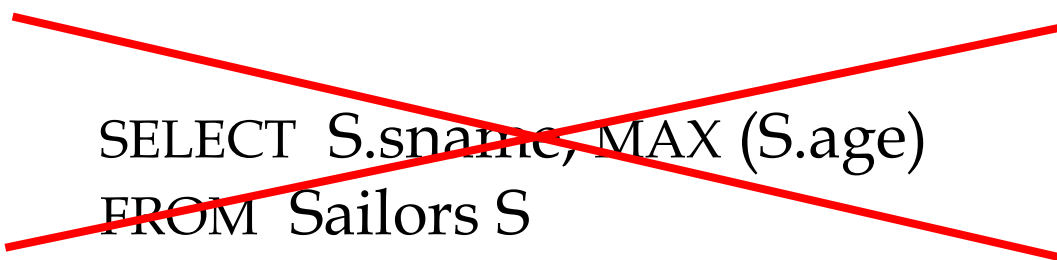
Find the names of the sailors with the highest rating.

```
SELECT S.sname
FROM Sailors S
WHERE S.rating=
      (SELECT MAX(S2.rating)
       FROM Sailors S2)
```

sid	sname	rating	age
1	Fred	7	20
2	Jim	2	39
9	Mike	7	20
4	Mary	1	17
22	Fred	7	50
3	Nancy	2	21

Find name and age of the oldest sailor(s)

- The first query is illegal: If the SELECT clause uses an aggregate operation, then it must use *only* aggregate operations unless the query contains **GROUP BY** clause)



```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
      (SELECT MAX (S2.age)  
       FROM Sailors S2)
```

GROUP BY and HAVING

- So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

**Find the age of the youngest sailor for each
rating level**

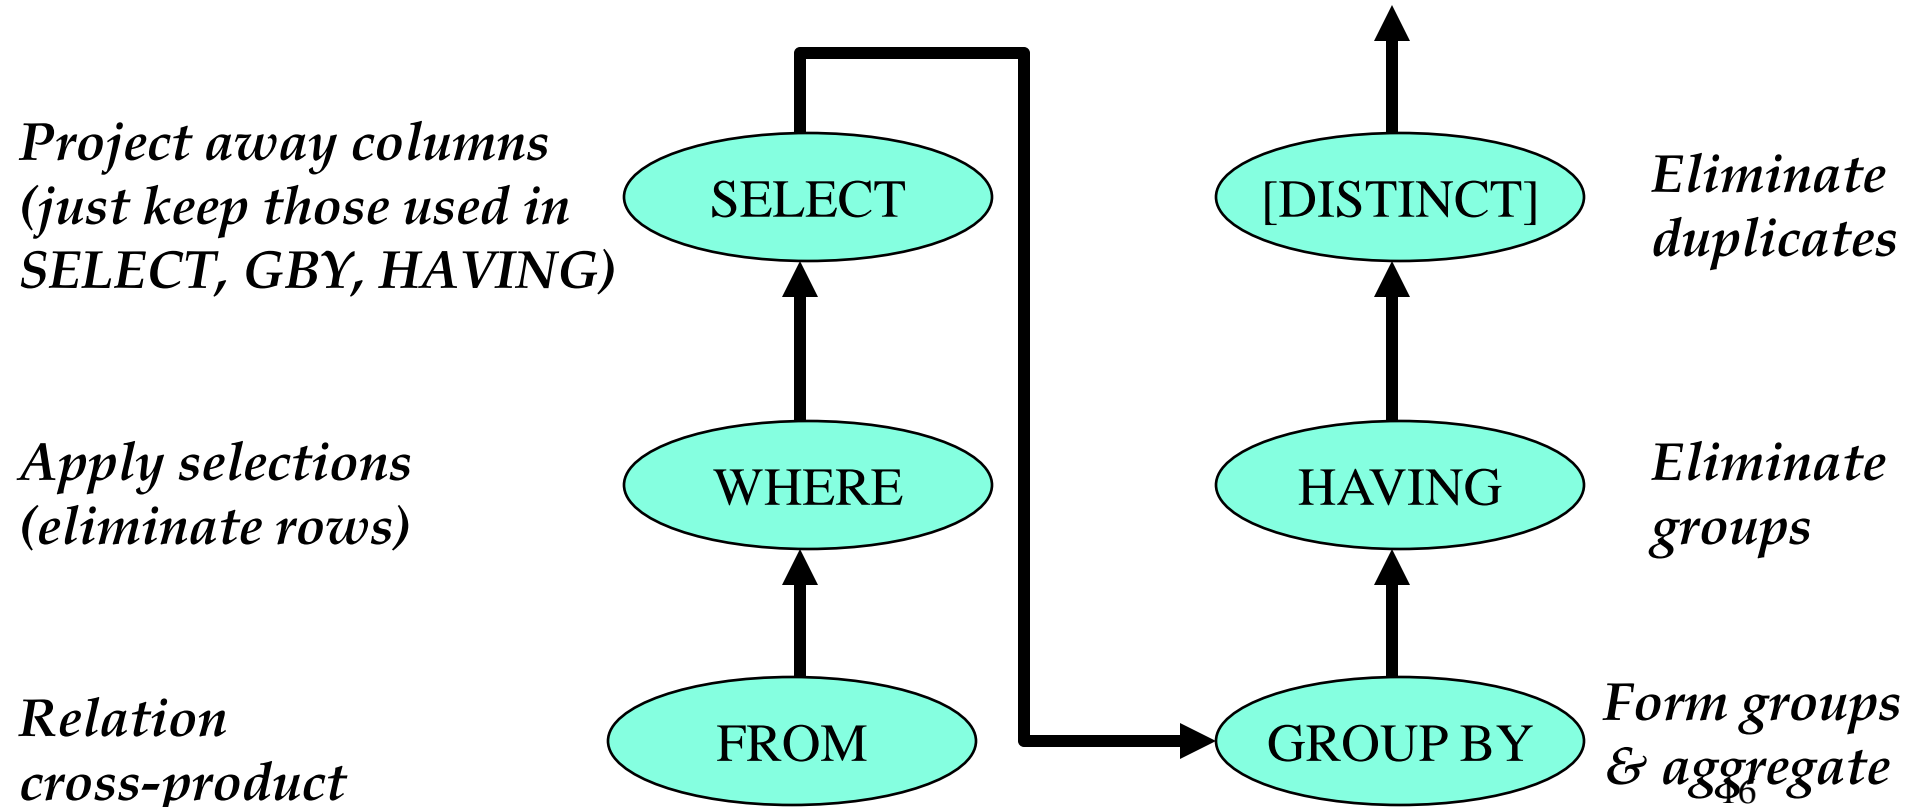
```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating;
```

Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) >1;
```

Conceptual Evaluation

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>



Find the age of the youngest sailor for each rating level with at least 2 sailors

```
SELECT S.rating, MIN (S.age)
```

```
FROM Sailors S
```

```
GROUP BY S.rating
```

```
HAVING S.age > 30
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

sid sname rating age

~~29 brutus 1 33.0~~

22 dustin 7 45.0

64 horatio 7 35.0

~~31 lubber 8 55.5~~

58 rusty 10 35.0

71 zorba 10 16.0

Expressions in *group-qualification* must have a single value per group!

One answer tuple is generated **per** qualifying group.

Find the age of the youngest sailor **for each rating level with at least 2 sailors**

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) >1
```

Equivalently...

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
            FROM Sailors S2
            WHERE S.rating=S2.rating)
```

sid	sname	rating	age
-----	-------	--------	-----

29	brutus	1	33.0
----	--------	---	------

22	dustin	7	45.0
----	--------	---	------

64	horatio	7	35.0
----	---------	---	------

31	lubber	8	55.5
----	--------	---	------

58	rusty	10	35.0
----	-------	----	------

71	zorba	10	16.0
----	-------	----	------

- Shows HAVING clause can also contain a subquery.

Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

The diagram shows two subset relationships indicated by red arrows and the symbol \subseteq . One arrow points from *qualification* to *grouping-list* with the symbol \subseteq to its right. Another arrow points from *group-qualification* to *grouping-list* with the symbol \subseteq to its right.

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - The attribute list (i) must be **a subset of** *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have **a single value per group**. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

Conceptual Evaluation

- The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a single value per group!
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*.
- One answer tuple is generated **per qualifying group**.

For each red boat, find the number of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

- Grouping over a join of two relations.

**For each red boat, find the number of
reservations for this boat**

Reserves

sid	bid	date
1	101	1/1/2017
1	102	15/1/2017
2	108	3/3/2016
1	101	5/6/2016
1	108	4/4/2017

Boats

bid	color
101	red
102	green
108	red



sid	R.bid	date	B.bid	color
1	101	...	101	red
1	102		102	green
2	108		108	red
1	101		101	red
1	108		108	red



sid R.bid B.bid color

1	101	...	101	red
1	101	...	101	red
2	108	...	108	red
1	108	...	108	red

For each red boat, find the number of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

- Grouping over a join of two relations.
- What do we get if we remove *B.color='red'* from the WHERE clause and add a HAVING clause with this condition?

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
FROM Sailors S2
WHERE S.rating=S2.rating)
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

For rating=7
Count(*)
2

For rating=10
Count(*)
2

rating	
7	35.0
10	35.0

- Shows HAVING clause can also contain a subquery.