

SQL: Structured Query Language

Part IV

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 sailors (of any age)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
FROM Sailors S2
WHERE S.rating=S2.rating)
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

For rating=7
Count(*)
2

For rating=10
Count(*)
2

rating	
7	35.0
10	35.0

- Shows HAVING clause can also contain a subquery.

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

- Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes '*unnecessary*'.

Rating	age
7	45.0
8	55.5
7	35.0
1	33.0
10	35.0

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Rating	age
1	33.0
7	45.5
7	35.0
8	55.5
10	35.0

rating	
7	35.0

Answer relation

Find the age of the youngest sailor with age \geq 18, for each rating with at least 2 such sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

Find the age of the youngest sailor with age \geq 18, for each rating with at least 2 such sailors

SELECT Temp.rating, Temp.minage
FROM



WHERE Temp.rcount > 1 Assume a TEMP storing number
 & min age of sailors \geq 18 per rating

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0



rating	rcount	minage
1	1	33.0
7	2	35.0
8	1	55.5
10	1	35.0

Find the age of the youngest sailor with age \geq 18, for each rating with at least 2 such sailors

```
SELECT Temp.rating, Temp.minage  
FROM
```

```
(SELECT S.rating, COUNT(*) AS rcount, MIN(S.age) AS minage  
FROM Sailors S
```

```
WHERE S.age  $\geq$  18
```

```
GROUP BY S.rating) AS Temp
```

```
WHERE Temp.rcount > 1
```

Assume a TEMP storing number
& min age of sailors \geq 18 per rating

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0



rating	rcount	minage
1	1	33.0
7	2	35.0
8	1	55.5
10	1	35.0

RECALL: Find name and age of the oldest sailor(s)

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

Find those ratings for which the average age is the minimum over all ratings

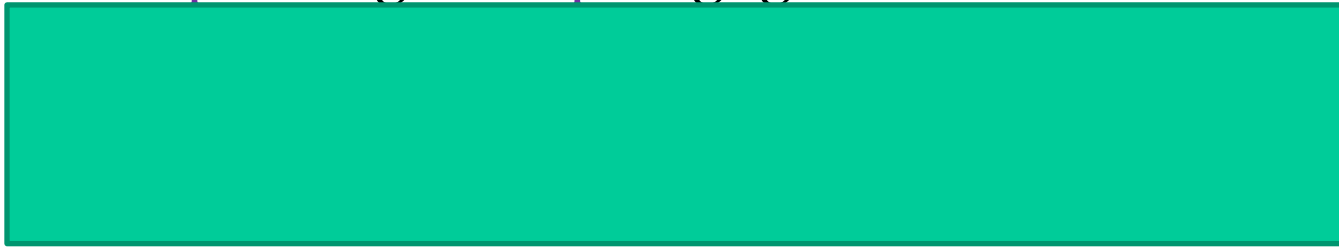
- Aggregate operations **cannot be nested!**
- The following is **WRONG:**

```
SELECT S.rating  
FROM Sailors S  
WHERE S.age = (SELECT MIN (AVG (S2.age)) FROM Sailors S2)
```


Find those ratings for which the average age is the minimum over all ratings

➤ Correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage  
FROM
```



```
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)  
FROM Temp)
```

OK to define once
on pen&paper

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

Find those ratings for which the average age is the minimum over all ratings

➤ Correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp)
```

OK to define once
on pen&paper

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	avgage
1	33.0
7	40.0
8	55.5
10	25.5

Find those ratings for which the average age is the minimum over all ratings

- How about....?

~~SELECT Temp.rating, MIN(Temp.avgage)
FROM (SELECT S.rating, AVG (S.age) AS avgage
FROM Sailors S
GROUP BY S.rating) AS Temp
GROUP BY Temp.rating~~

Temp	rating	avgage
	1	33.0
	7	40.0
	8	55.5
	10	25.5

If the SELECT clause uses an aggregate operation, then it must use *only* aggregate operations unless the query contains **GROUP BY** clause


~~SELECT S.sname, MAX (S.age)
FROM Sailors S~~

Find those ratings for which the average age is the minimum over all ratings

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp)
```

What if you can't use MIN?

What if you can't use TEMP?



rating	avgage
--------	--------

1	33.0
---	------

7	40.0
---	------


8	55.5
---	------

10	25.5
----	------

Find those ratings for which the average age is the minimum over all ratings

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp)
```

What if you can't use MIN?



```
SELECT S.rating, AVG (S.age) AS avgage
FROM Sailors S
GROUP BY S.rating
HAVING AVG (S.age) <= ALL
```

What if you can't use TEMP?

rating avgage

1 33.0

7 40.0

8 55.5

10 25.5

```
(SELECT AVG(S2.age)
FROM Sailors S2
GROUP BY S2.rating)
```

33.0

40.0

55.5

25.5

Null Values

- Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- The presence of *null* complicates many issues.
 - Special operators “IS NULL” and “IS NOT NULL”.

Null Values: 3 valued logic

- Is $rating > 8$ true or false when $rating$ is equal to *null*? What about **AND**, **OR** and **NOT** connectives?
- We need a 3-valued logic (true, false and *unknown*).

sid	sname	rating	age
21	Dan	NULL	38

- $rating > 8$ can be **true**, **false** or **unknown**
- $rating > 8$ OR $age < 40$???
- $rating > 8$ OR $age > 40$???
- $rating > 8$ AND $age > 40$???

Null Values: 3 valued logic

- Is $rating > 8$ true or false when $rating$ is equal to *null*? What about **AND**, **OR** and **NOT** connectives?
- We need a 3-valued logic (true, false and *unknown*).

sid	sname	rating	age
21	Dan	NULL	38

- $rating > 8$ can be **true**, **false** or **unknown**
- $rating > 8$ OR $age < 40$???
- $rating > 8$ OR $age > 40$???
- $rating > 8$ AND $age > 40$???

OR TABLE		
T	T	T
T	F	T
F	T	T
T	Unknown	T
Unknown	T	T
F	F	F
F	Unknown	Unknown
Unknown	F	Unknown
Unknown	Unknown	Unknown

Meaning of constructs must be defined carefully. (e.g., WHERE clause **eliminates rows** that **don't evaluate to true**.)

Null Values: Aggregate Operators

- COUNT(*) → counts also the null values
- All others (as below): **discards** NULL values

COUNT ([DISTINCT] A)

SUM ([DISTINCT] A)

AVG ([DISTINCT] A)

MAX (A)

MIN (A)

Outer Joins


- Left outer join for S and R:
 - Each S row without a matching R rows appears (once) in the result, with R columns including NULLs
 - SELECT S.sid, R.bid
- FROM Sailors S **NATURAL LEFT OUTER JOIN** Reserves R

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0

Reserves

sid	bid	date
22	101	...
31	102	



sid	bid
22	101
31	102
71	null

Modifying Tables – Insert

- Inserting a single row into a table
 - Attribute list can be omitted if it is the same as in CREATE TABLE
 - NULL and DEFAULT values can be specified

```
INSERT INTO Transcript(StudId, CrsCode, Semester, Grade)  
VALUES (12345, 'CSE305', 'S2000', NULL)
```

Bulk Insertion

- Insert the rows output by a SELECT

```
CREATE TABLE DeansList (  
    StudId      INTEGER,  
    Credits     INTEGER,  
    CumGpa      REAL,  
    PRIMARY KEY (StudId))
```

```
INSERT INTO DeansList (StudId, Credits, CumGpa)  
SELECT      T.StudId, 3 * COUNT (*),  AVG(T.Grade)  
FROM        Transcript T  
GROUP BY    T.StudId  
HAVING      AVG (T.Grade) > 3.5  AND  
            COUNT(*) > 30
```

Recall: Transcript(*StudId*, *CrsCode*, *Semester*, *Grade*)

Modifying Tables – Delete

- Similar to SELECT except:
 - No project list in DELETE clause
 - **No Cartesian product** in FROM clause (only 1 table name)
 - Rows satisfying WHERE clause (general form, including subqueries, allowed) are deleted instead of output

```
DELETE FROM Transcript T
WHERE T.Grade IS NULL AND T.Semester <> 'F2020'
```

Modifying Data - Update

```
UPDATE Employee E  
SET      E.Salary = E.Salary * 1.05  
WHERE    E.Department = 'R&D'
```

- Updates rows in a **single table**
- All rows satisfying WHERE clause (general form, including subqueries, allowed) are updated

Test first!

```
UPDATE Employee E
SET      E.Salary = E.Salary * 5
WHERE    E.Salary <= (SELECT AVG(E2.salary)
                       FROM Employee E2)
```



eid	ename	salary	age
1	Fred	1000	20
2	Jim	2000	39
9	Mike	500	20
4	Mary	3000	17
22	Fred	500	50
3	Nancy	1400	21

Avg: $8400/6 = 1400$