

Hash-Based Indexing

Motivation

- The primary goal is to locate the desired record in a single access of disk.
 - Sequential search: $O(N)$ (N : no of blocks)
 - B+ trees: $O(\log_k N)$
 - Hashing: $O(1)$, a **single (one) disk-access** method
- In hashing, the search key of a record is transformed into an address and the record is stored at that address.
- Hash-based indexes are the best for **equality selections**. **Can not** support **range searches**.
- Static and dynamic hashing techniques exist.

Hash-based Index

- Data entries are kept in *buckets* (an abstract term)
- Each bucket is a collection of **one primary page** and **zero or more overflow pages**.
- Given a search key value, k , we can find the bucket where the **data entry k^*** is stored as follows:
 - Use a *hash function*, denoted by h
 - The value of $h(\text{key})$ is the address for the desired bucket.
 - $h(\text{key})$ should distribute the search key values *uniformly* over the collection of buckets

Hash Functions

- It is always a bad idea to use field values themselves for hashing, as real data is *almost never* uniformly distributed.
 - Hash function is for transforming a non-uniform distribution to a (nearly) uniform one.
- **Key mod N:** N is the **number of buckets**, better if it is prime.
- **Folding:** e.g. 123|456|789: add them and take mod.
- **Truncation:** e.g. 123456789 map to a table of 1000 addresses by picking 3 digits of the key.
- **Squaring:** Square the key and then truncate
- **Radix conversion:** e.g. 1 2 3 4 treat it to be base 11, truncate if necessary.
- Hash functions **do not preserve order!** → Hence, hash based indexes **do not help for range search.**

Static Hashing

- In a file organized using static hashing, the data entries k^* are stored in the buckets in the **primary area** and, possibly, in the **overflow area**.
- **Primary Area:** # primary pages fixed, **never de-allocated**; (say **M** buckets).
 - A hash function works on the *search key* field, and maps keys to values from **0** to **M-1**
 - A simple hash function: $h(key) = f(key) \bmod M$
- **Overflow area:** disjoint from the primary area. It keeps overflow pages which hold records whose key maps to a full bucket.
 - **Chaining:** The address of an overflow page is added to the **overflow chain** of the full bucket

Static Hashing

- **Bucket factor** (Bkfr) is the number of data entries that can be held at **a bucket**.
- *Collision* does not cause a problem as long as there is still room in the mapped bucket. *Overflow* occurs during insertion when a record is hashed to the bucket that is already full.

Example

- Assume $f(\text{key}) = \text{key}$. Let $M = 5$. So, $h(\text{key}) = \text{key} \bmod 5$
- Bucket factor (Bkfr) = 3 records.

Insert: 12, 35, 44, 60, 6, 46, 57, 33, 62, 17

Bkfr = 3

Buckets:
0 to M-1



0
1
2
3
4

35*	60*	
6*	46*	
12*	57*	62*
33*		
44*		



k^* are stored
(**data entries** using
Alt-1-2 or 3)



17*		
-----	--	--

Overflow area
(disjoint from
primary area)

Primary area

If no overflow, fetching a record costs only one disk access!

Load Factor (Packing density)

- To limit the amount of overflow we allocate more space to the primary area than we need (i.e. the primary area will be, say, 70% full)




- Load Factor =
$$\frac{\text{\# of records in the file}}{\text{\# of spaces in primary area}}$$

$$\Rightarrow Lf = \frac{n}{M * Bkfr}$$

Example: What is Load Factor?

$$LF = \frac{n}{M * Bkfr} = \frac{\text{9+1= 10 records}}{\text{5 * 3 cells}} = 2/3$$



Primary area
(primary bucket pages)

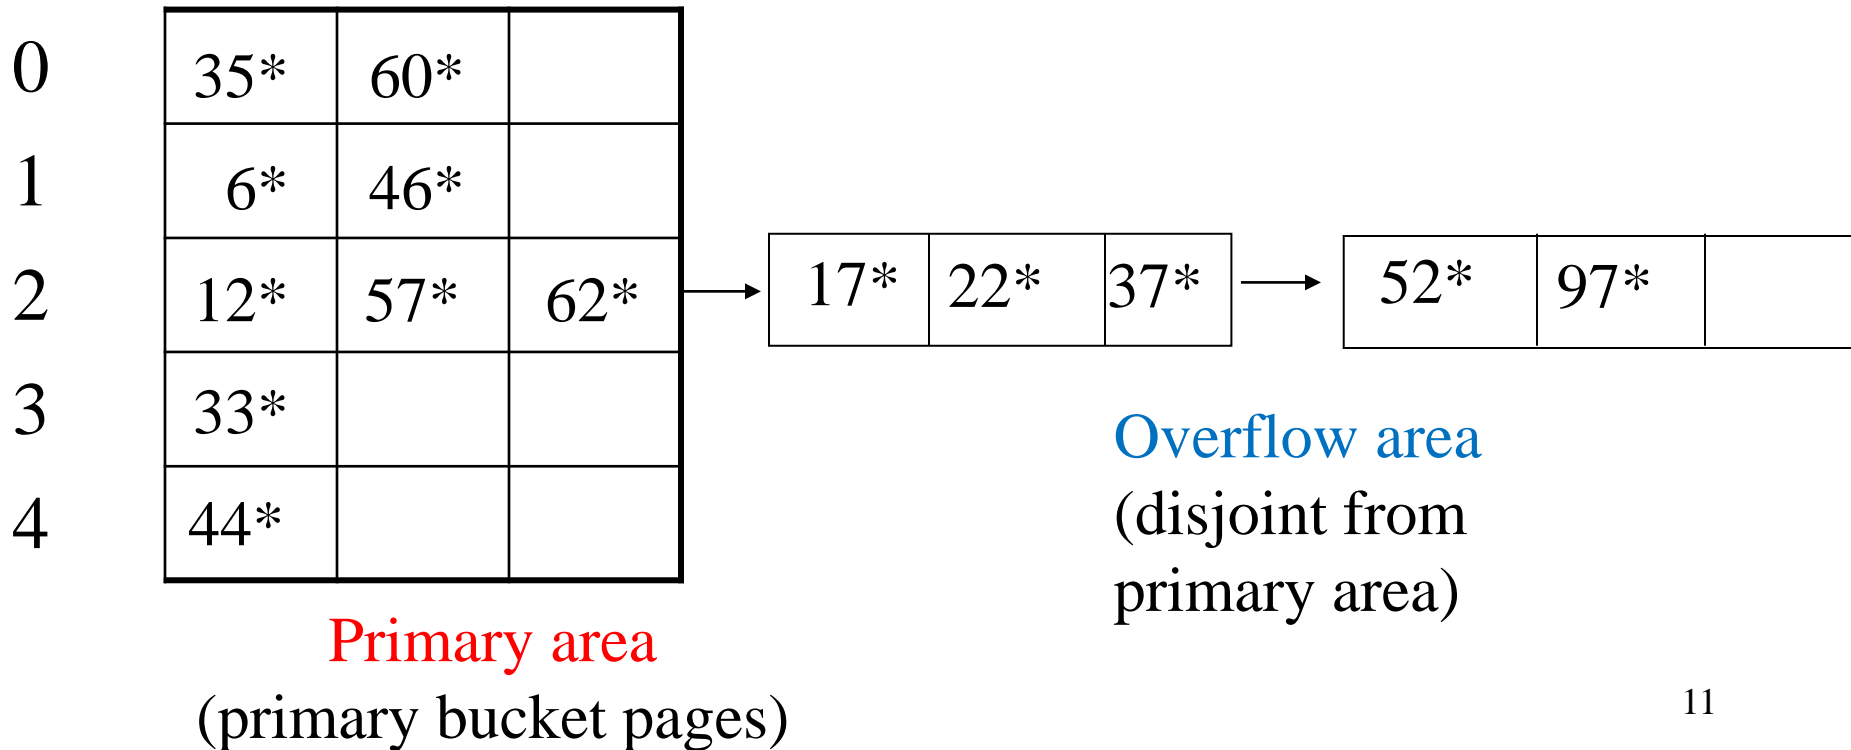
Overflow area
(disjoint from
primary area)

Effects of Lf and Bkfr

- Performance can be enhanced by the choice of Bkfr and load factor.
- In general, a smaller load factor means
 - less overflow and a faster fetch time;
 - but more wasted space.
- A larger Bkfr means
 - less overflow in general,
 - but slower fetch.

Insertion and Deletion

- Insertion: New records are inserted at the end of the chain.
- Insert: 22, 37, 52, 97

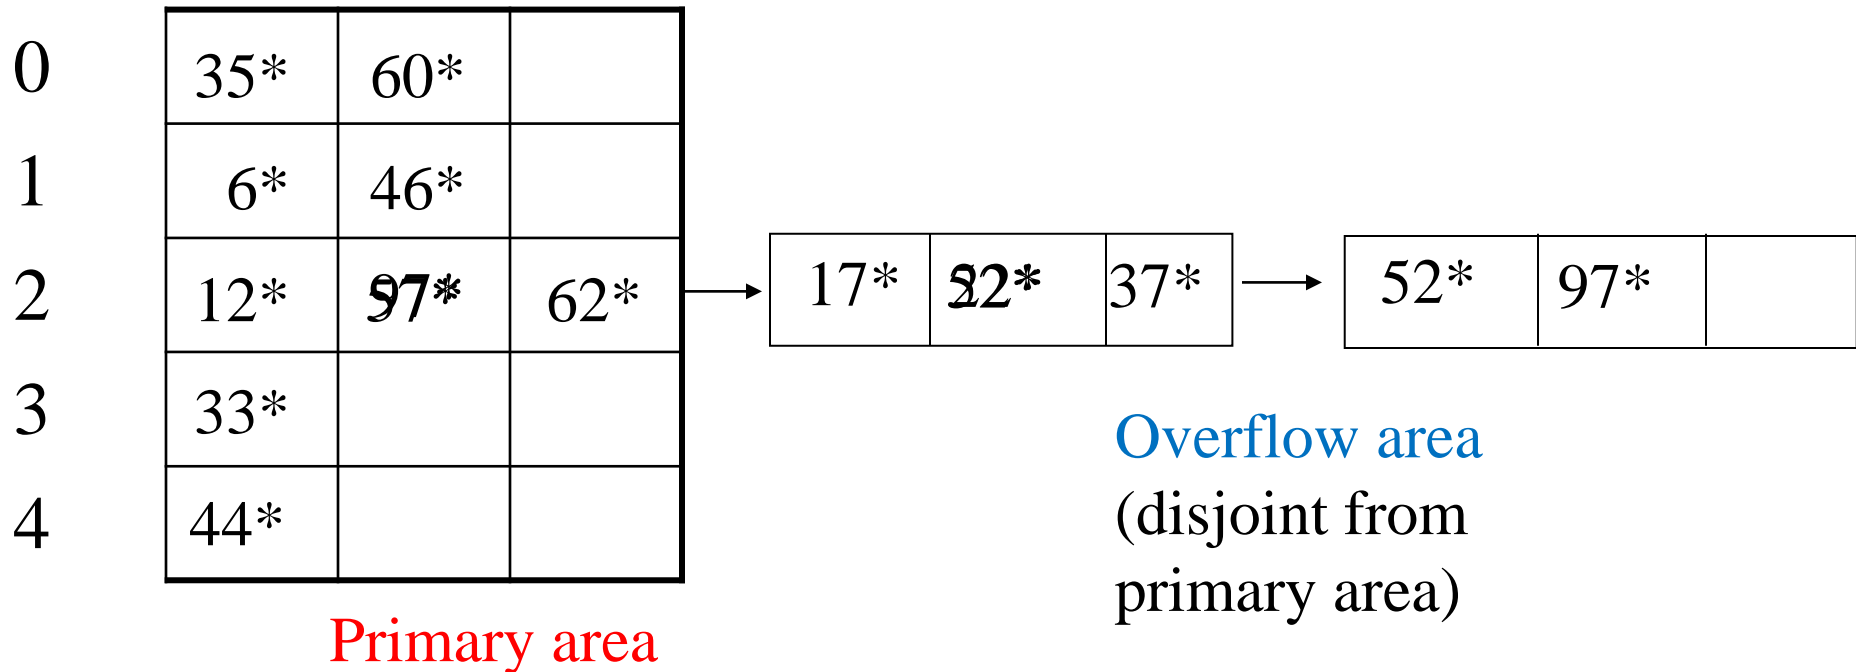


Insertion and Deletion

- Deletion: Two ways are possible:
 1. Mark the record to be deleted
 2. Consolidate sparse buckets when deleting records.
- In the 2nd approach:
 - When a record is deleted, fill its place with the last record in the chain of the current bucket.
 - Deallocate the last bucket when it becomes empty.

Deletion

- When a record is deleted, fill its place with the last record in the chain of the current bucket.
- Deallocate the last bucket when it becomes empty.
- Delete: 57, then 22



Problem of Static Hashing

- The main problem with **static** hashing: **the number of buckets is fixed**:
 - **Long overflow chains** can develop and degrade performance.
→ would require **re-organization** at some point
 - On the other hand, if a file shrinks greatly, a lot of bucket space will be wasted.
- There are some other hashing techniques that allow **dynamically growing and shrinking** hash index. These include:
 - extendible hashing
 - linear hashing


Extendible Hashing

Extendible Hashing

- **Basic Idea:**
 - No overflow buckets
 - Instead add a level of indirection
- Use **directory of pointers** to buckets
- Double # of buckets by doubling the directory
 - Directory much smaller than file, so doubling it is much cheaper.
- Split only the bucket that just overflowed!
 - Adjust the hash function


General Structure

- Directory is an array of size 4, so 2 bits needed.
- Bucket for record r with key k is in the element with index = '*global depth*' least significant bits of $h(k)$;



Search

- To search for a data entry, apply a hash function h to the key and take **the last d bits** of its binary representation to reach the directory entry, which points the required bucket
- Example: search for 5^*




Handling Inserts

- Find bucket where record belongs.
- If there's room, put it there.
- Else, if bucket is full, split it:
 - increment **local depth** of original page
 - allocate new page with new **local depth**
 - re-distribute records from original page
 - add entry for the new page to the directory
 - double the directory if necessary

Example: Insert 21*, 19*, 15*

Assume $h(\text{key}) = \text{key}$ (in binary)

- $21 = 10101$
- $19 = 10011$
- $15 = 01111$



Example

if bucket is full, split it:

increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.

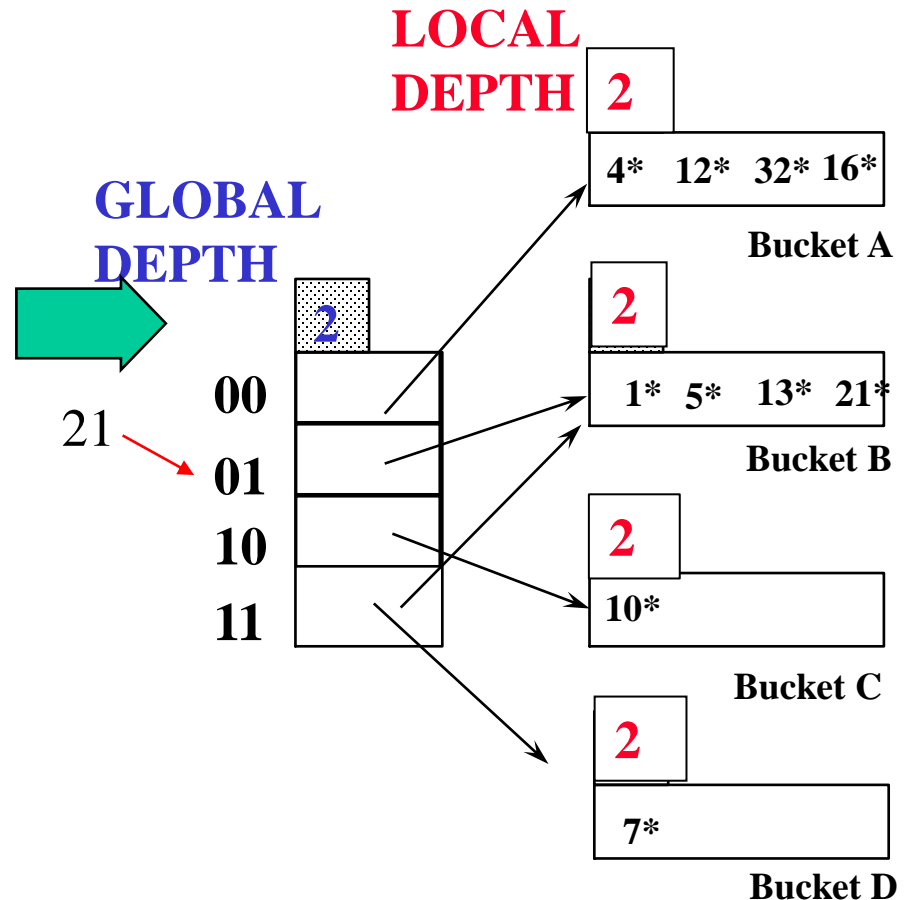
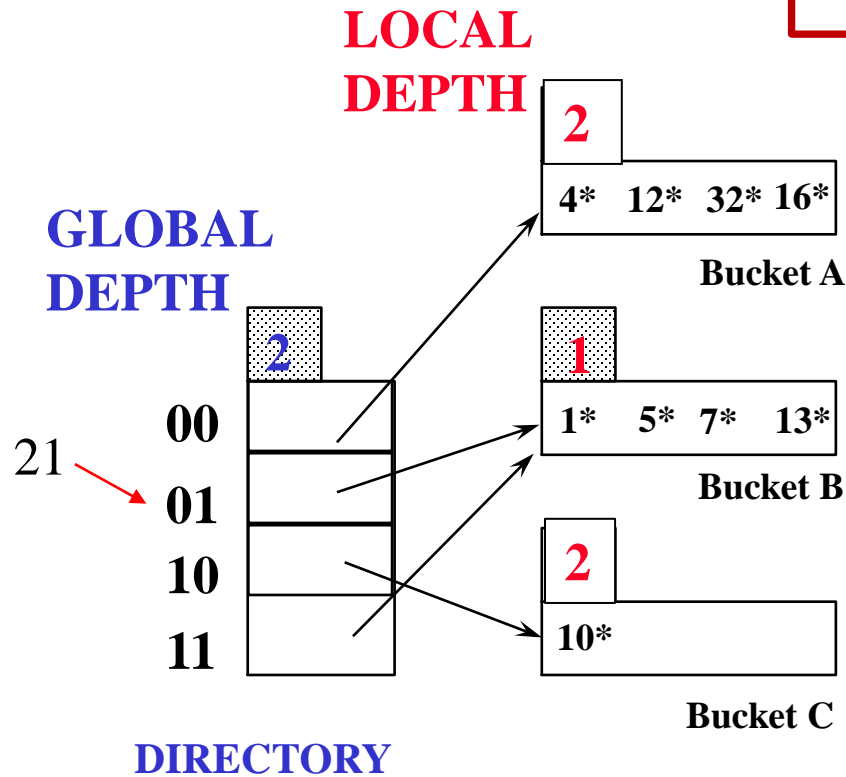
add entry for the new page to the directory

double the directory *if necessary*

➡ 21 = 10101

• 19 = 10011

• 15 = 01111




Example: Insert 21*, 19*, 15*

- 21 = 10101

➔ 19 = 10011

- 15 = 01111

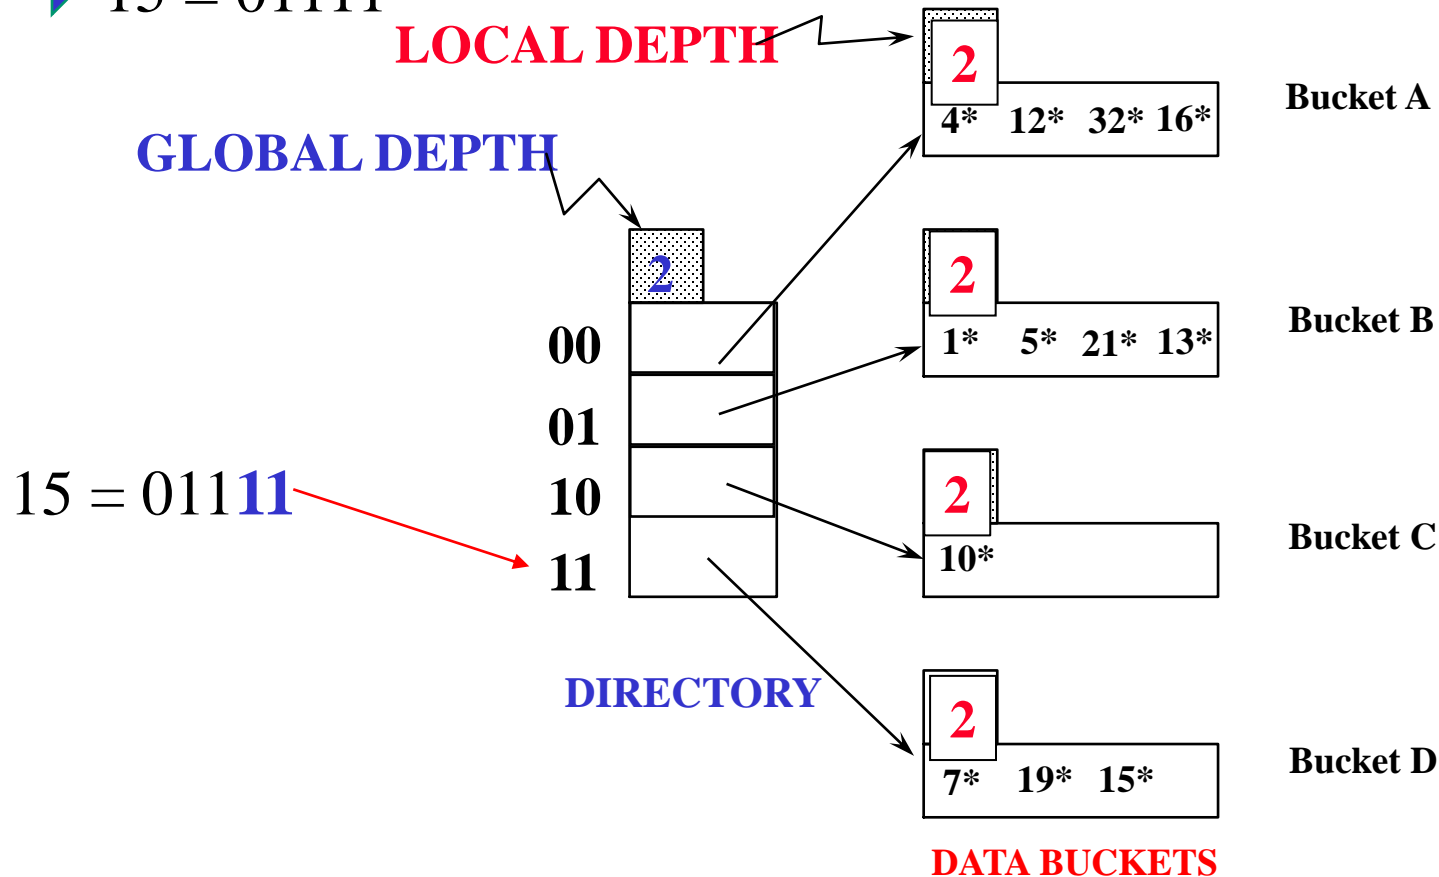


Example: Insert 21*, 19*, 15*

- 21 = 10101

- 19 = 10011

➡ 15 = 01111



Example:

Now, insert 20*

- insert 20*

if bucket is full, split it:

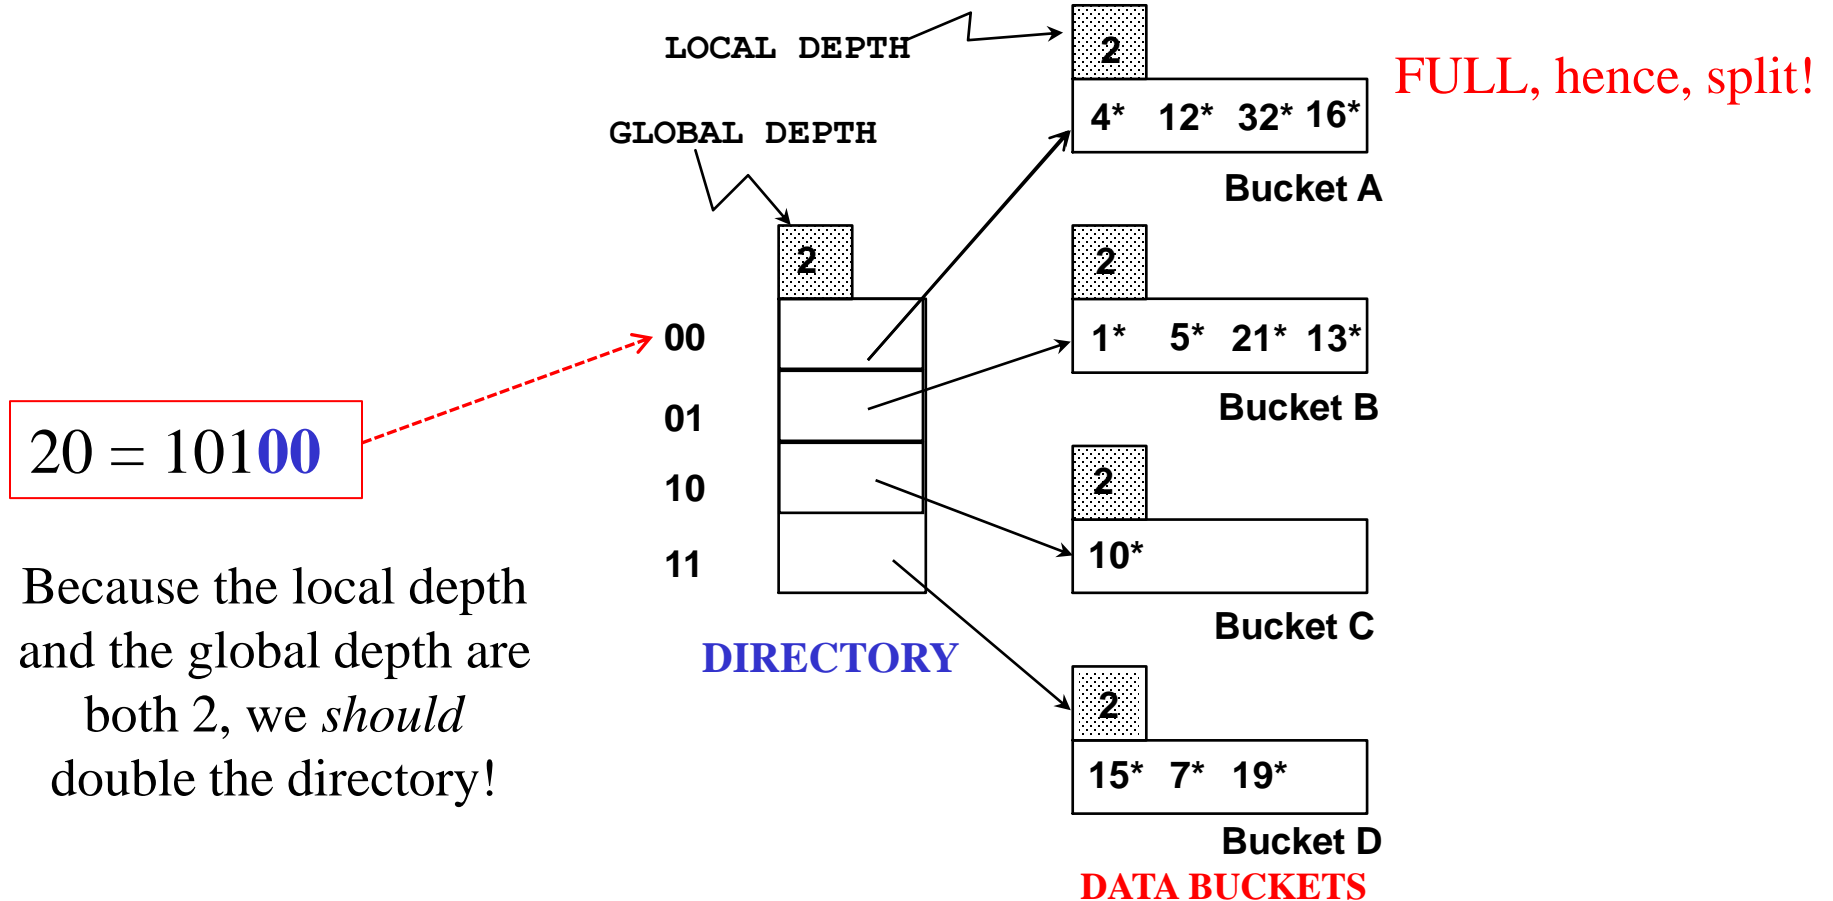
increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.

add entry for the new page to the directory

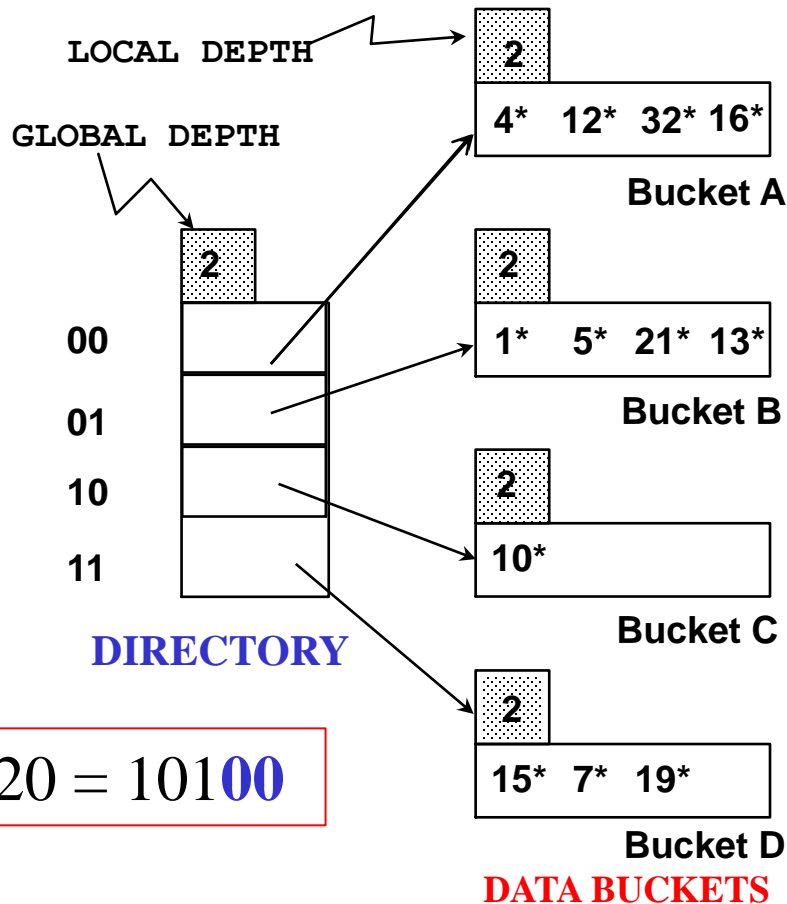
double the directory *if necessary*



Example:

Insert now 20*

- insert 20*



if bucket is full, split it:

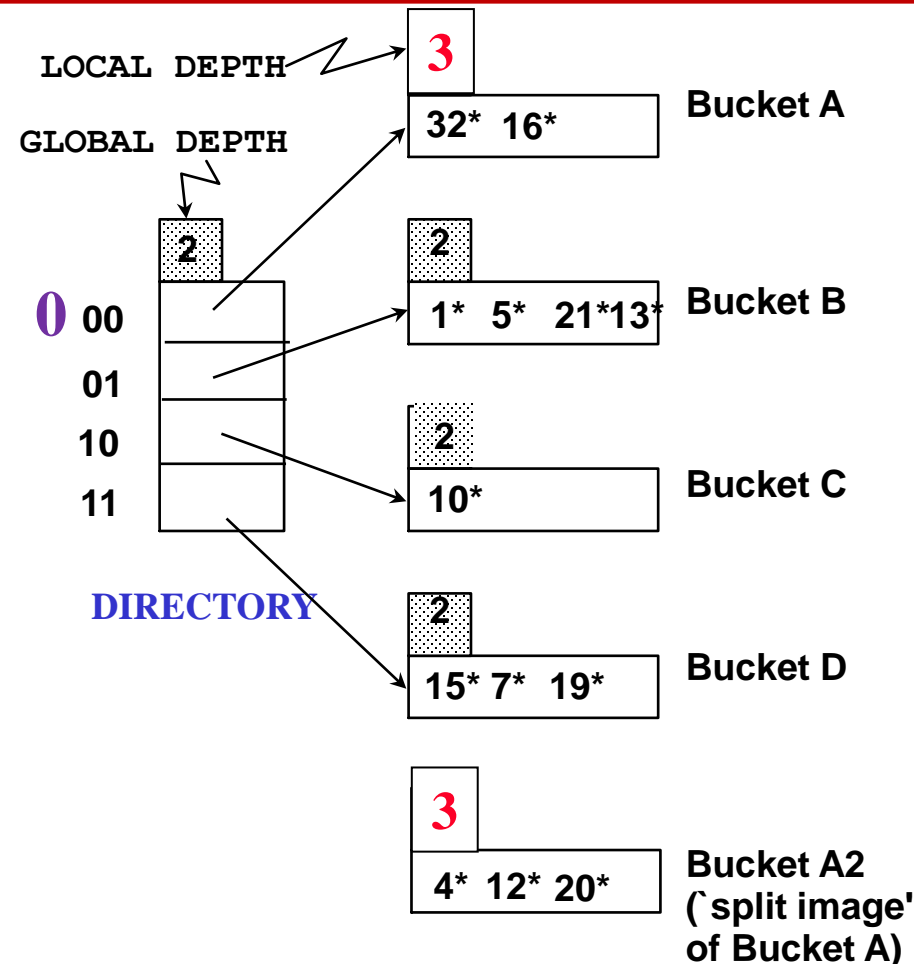
increment **local depth**

allocate new page with new **local depth**


re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary*

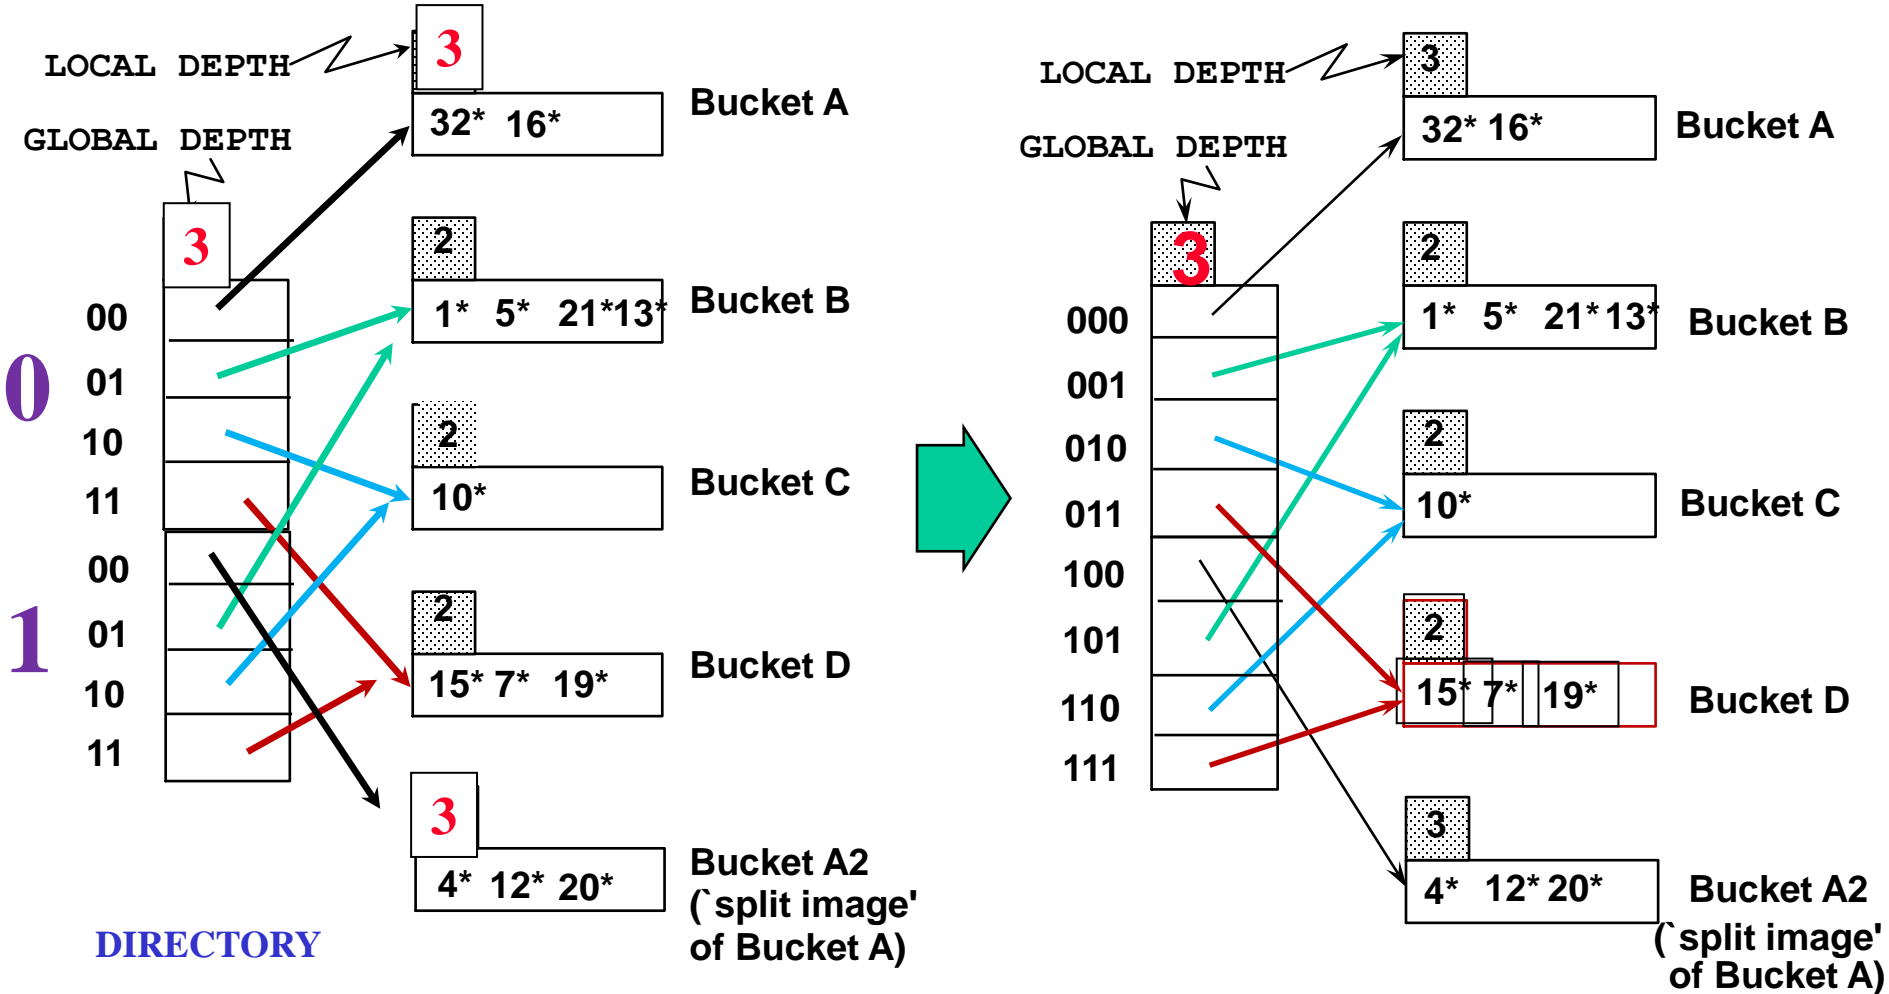


Doubling the directory (also increment GlobalDepth)



Example: Insert now 20*

Doubling the directory (also increment GlobalDepth)



Example:

Now insert 9*

- insert 9*

if bucket is full, split it:

increment **local depth**

allocate new page with new **local depth**

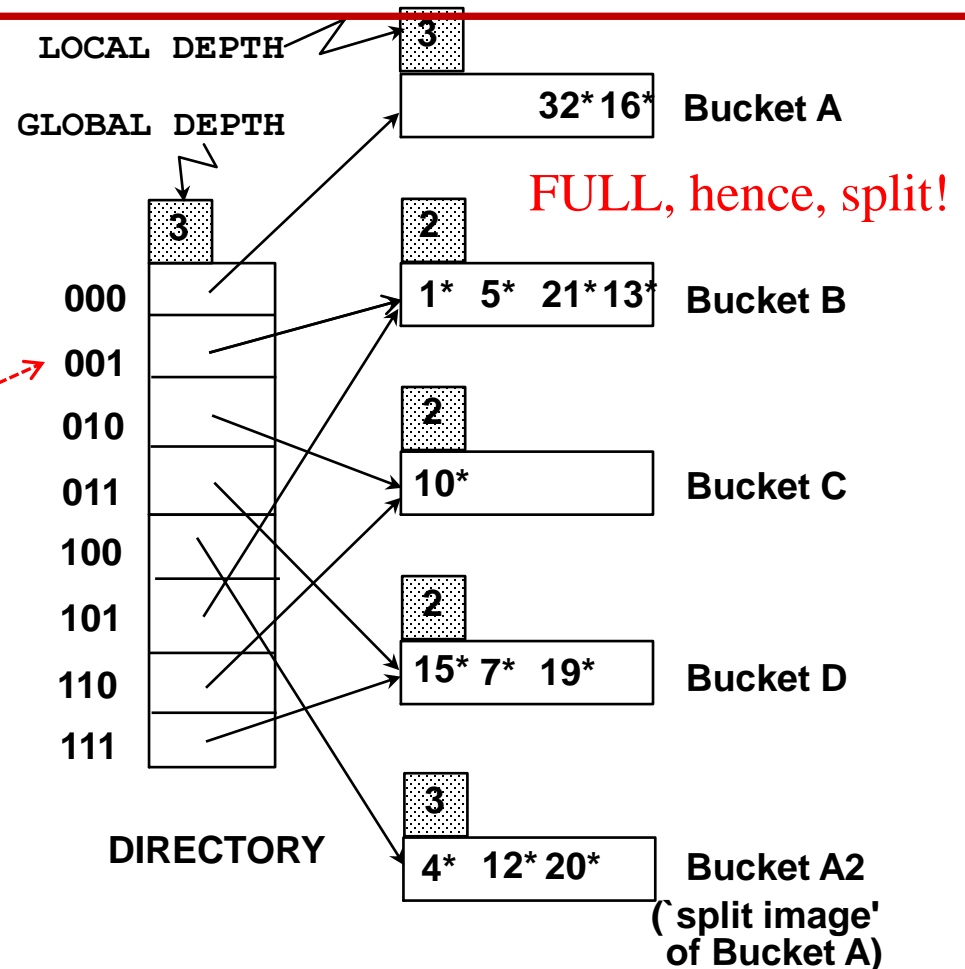
re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary* \rightarrow if $ORG\ LD = GD$

9 = 1001

Because the local depth
(i.e., 2) is *less than* the
global depth (i.e., 3),
NO need to double the
directory



Example:

Insert now 9*

if bucket is full, split it:

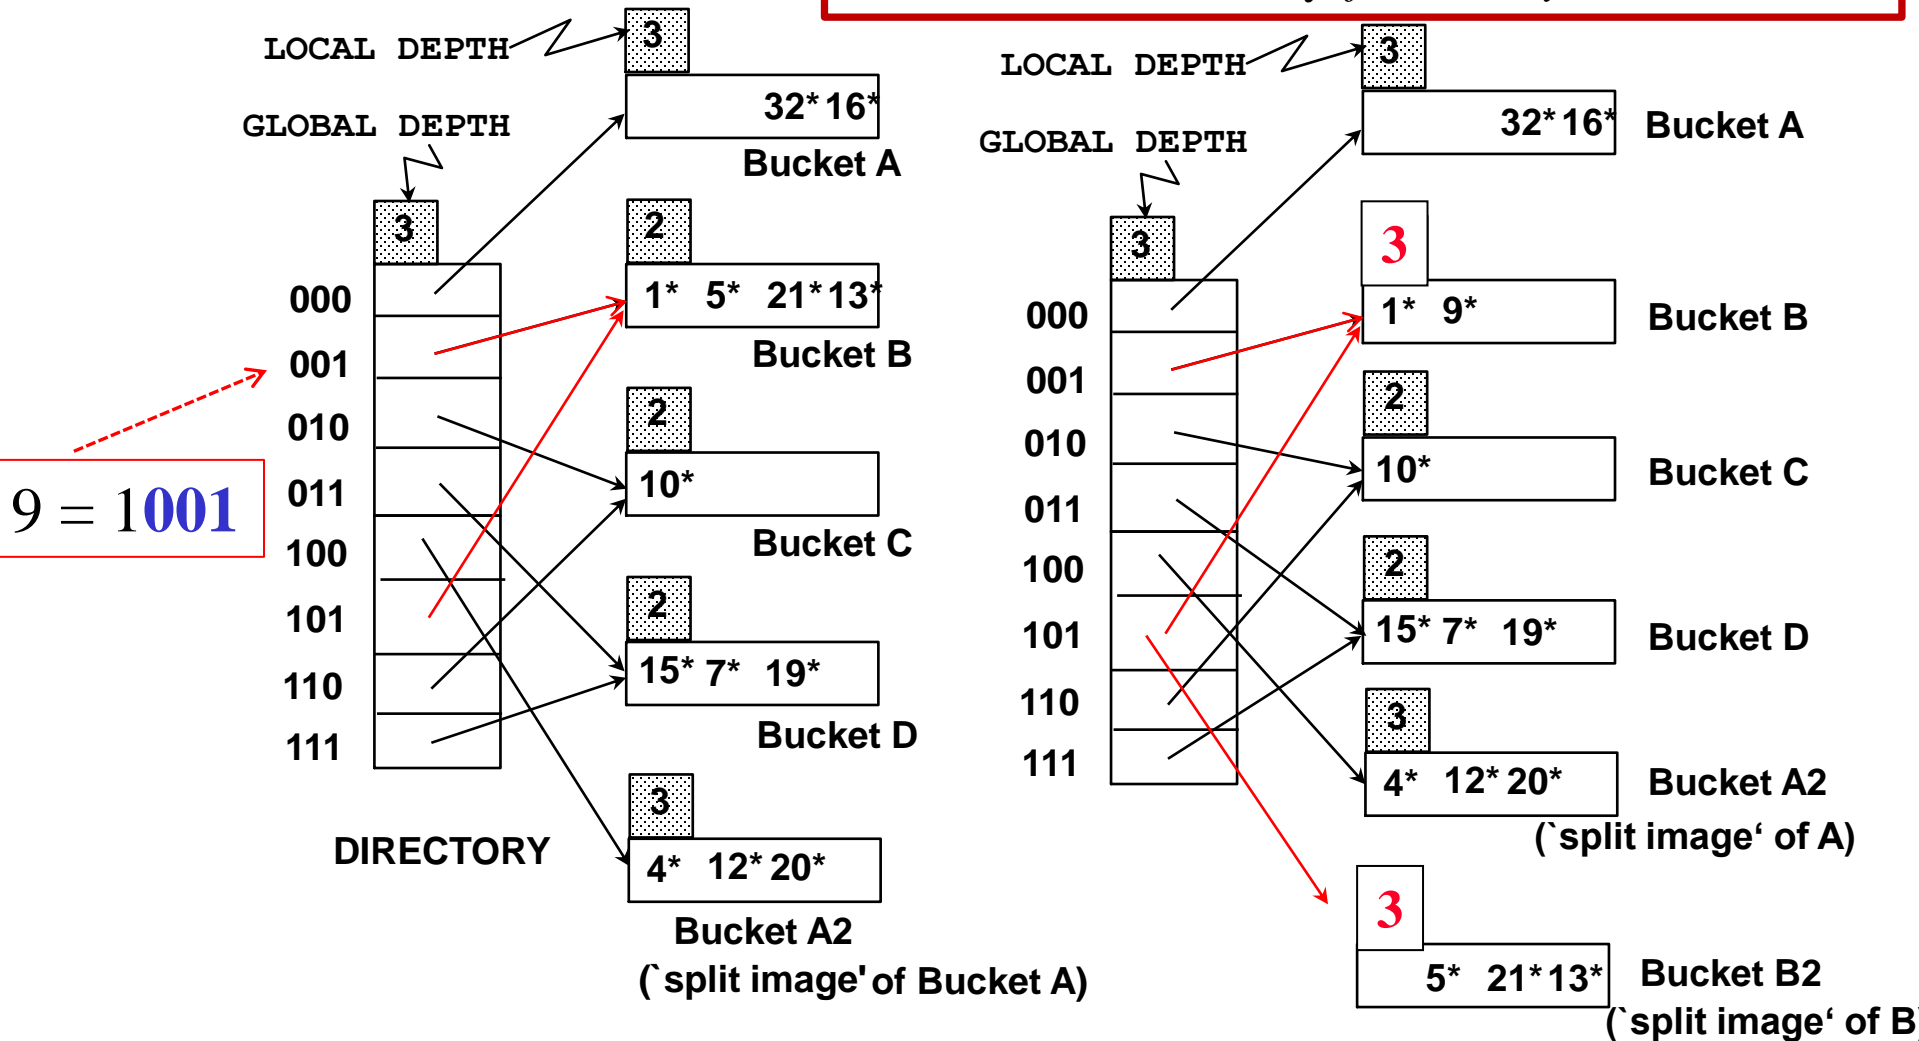
increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary*

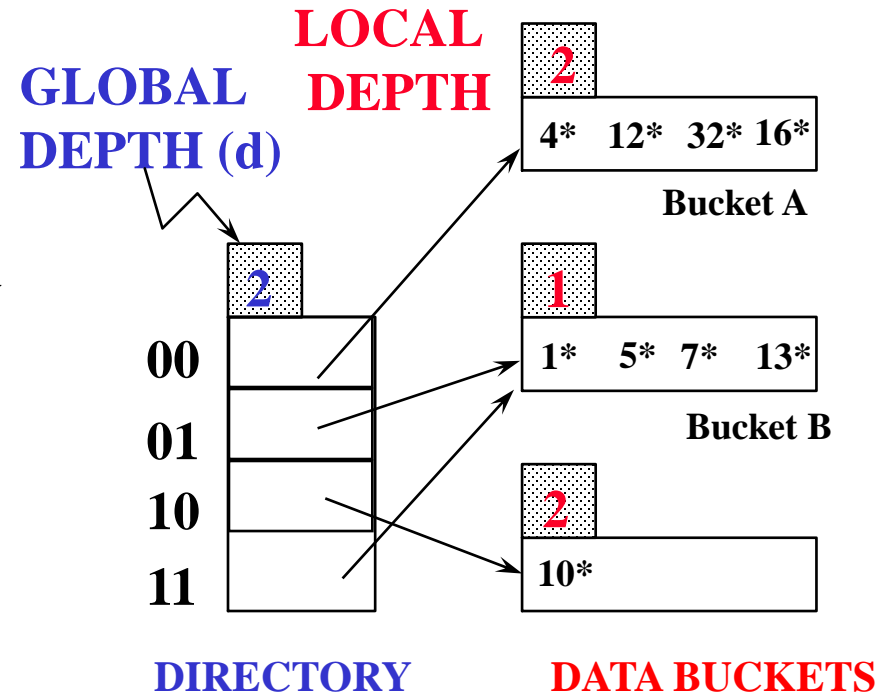


Points to Note

- *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
- *Local depth of a bucket*: # of bits used to determine if an entry belongs to this bucket.
- When does split cause directory doubling?
 - Before insert, *local depth* of bucket = *global depth*. Insert causes *local depth* to become $>$ *global depth*; directory is doubled by *copying it over* and 'fixing' pointer to split image page.

Extendible Hashing

- **Basic Idea:**
 - No overflow buckets
 - A level of indirection: a **directory of pointers** to buckets
- Double the directory periodically
 - Directory much smaller than file, so doubling is cheaper.
- Split only the bucket that just overflowed!
 - Adjust the hash function



Comments on Extendible Hashing

- If directory fits in memory, **equality search** answered with **one disk access**.
 - A typical example: a 100MB file with 100 bytes/entry and a page size of 4K contains 1,000,000 records (as data entries) but only about 25,000 directory elements
⇒ chances are high that directory will fit in memory.
- If the distribution *of hash values* is skewed (e.g., a large number of search key values all are hashed to the same bucket), directory can grow large.
 - But this kind of skew can be avoided with a well-tuned hashing function

Skewed Insertions

Add the following entries in sequence in an initially empty extensible hash file (**Bucketing factor is 2**)

$$16 = 10000$$

$$32 = 100000$$

$$4 = 100$$


(assume $h(k) = k$ in binary)

Skewed Insertions

Insert 16* and 32*

16 = 10000 and

32 = 100000



if bucket is full, split it:

increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.



add entry for the new page to the directory

double the directory *if necessary*

Skewed Insertions

Insert 4*

4 = 100



We need to double the directory

if bucket is full, split it:

increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.


add entry for the new page to the directory

double the directory *if necessary*


Skewed Insertions

Insert 4*

4 = 100



if bucket is full, split it:
increment **local depth**
allocate new page with new **local depth**
re-distribute records of original page.
add entry for the new page to the directory
double the directory *if necessary*



We need to double the directory again !

Skewed Insertions

Still inserting 4*

if bucket is full, *split* it:



increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary*



We need to double the directory again !

How many directory elements pointing to a bucket?



2^{GD-LD}

Deletion

- Locate data entry in its bucket and remove it.
- If removal of data entry makes bucket empty, can be merged with `split image`
- If each directory element points to same bucket as its split image, can halve directory.
 - Note: decreasing directory size is an expensive operation and should be done only if number of buckets becomes much smaller than the size of the directory

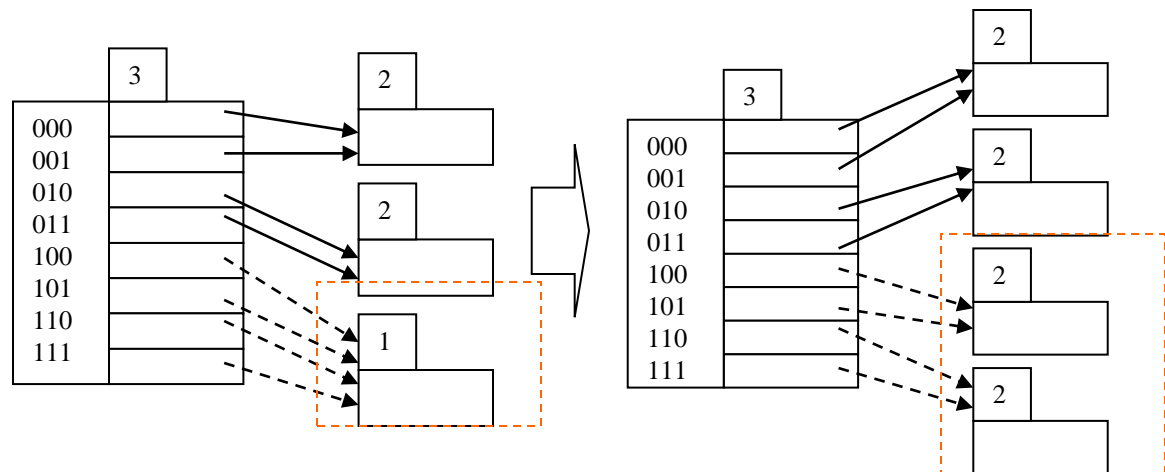
Can we use the most significant bits?

- Splitting (**Case 1: $i_j=i$**)
 - Only one element in directory (bucket address table) points to data bucket j
 - $i++$; split data bucket j to j, z ; $i_j=i_z=i$; rehash all items previously in j ;



Can we use the most significant bits ?


- Splitting (**Case 2: $i_j < i$**)
 - More than one element in directory (bucket address table) point to data bucket j
 - split data bucket j to j, z ; $i_j = i_z = i_j + 1$; Adjust the pointers previously point to j to j and z ; rehash all items previously in j ;



Directory Doubling

Why prefer least significant bits in directory
(instead of the *most* significant ones)?


Allows for doubling by copying the
directory and appending the new copy
to the original.



Skewed Insertions using most significant bits

Q2.(20 pts.) Consider an extendible hash structure. Each block can hold 2 records, *sorted* within each block in ascending order by their keys' hash values, and the structure is initially empty (i.e. a single entry directory of global depth 0 and a single data block (page) of local depth 0). We insert the following records in the given order using the leftmost k-bits of the hash value:

- (A) Key hashes to 00011 (= 3)
- (B) Key hashes to 10010 (= 18)
- (C) Key hashes to 10011 (= 19)
- (D) Key hashes to 00100 (= 4)
- (E) Key hashes to 10110 (= 22)
- (F) Key hashes to 11011 (= 27)
- (G) Key hashes to 10000 (= 16)



if bucket is full, *split* it:

increment **local depth**

allocate new page with new **local depth**

re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary*

- (A) Key hashes to 00011 (= 3)
- (B) Key hashes to 10010 (= 18)
- (C) Key hashes to 10011 (= 19)
- (D) Key hashes to 00100 (= 4)
- (E) Key hashes to 10110 (= 22)
- (F) Key hashes to 11011 (= 27)
- (G) Key hashes to 10000 (= 16)

if bucket is full, *split* it:

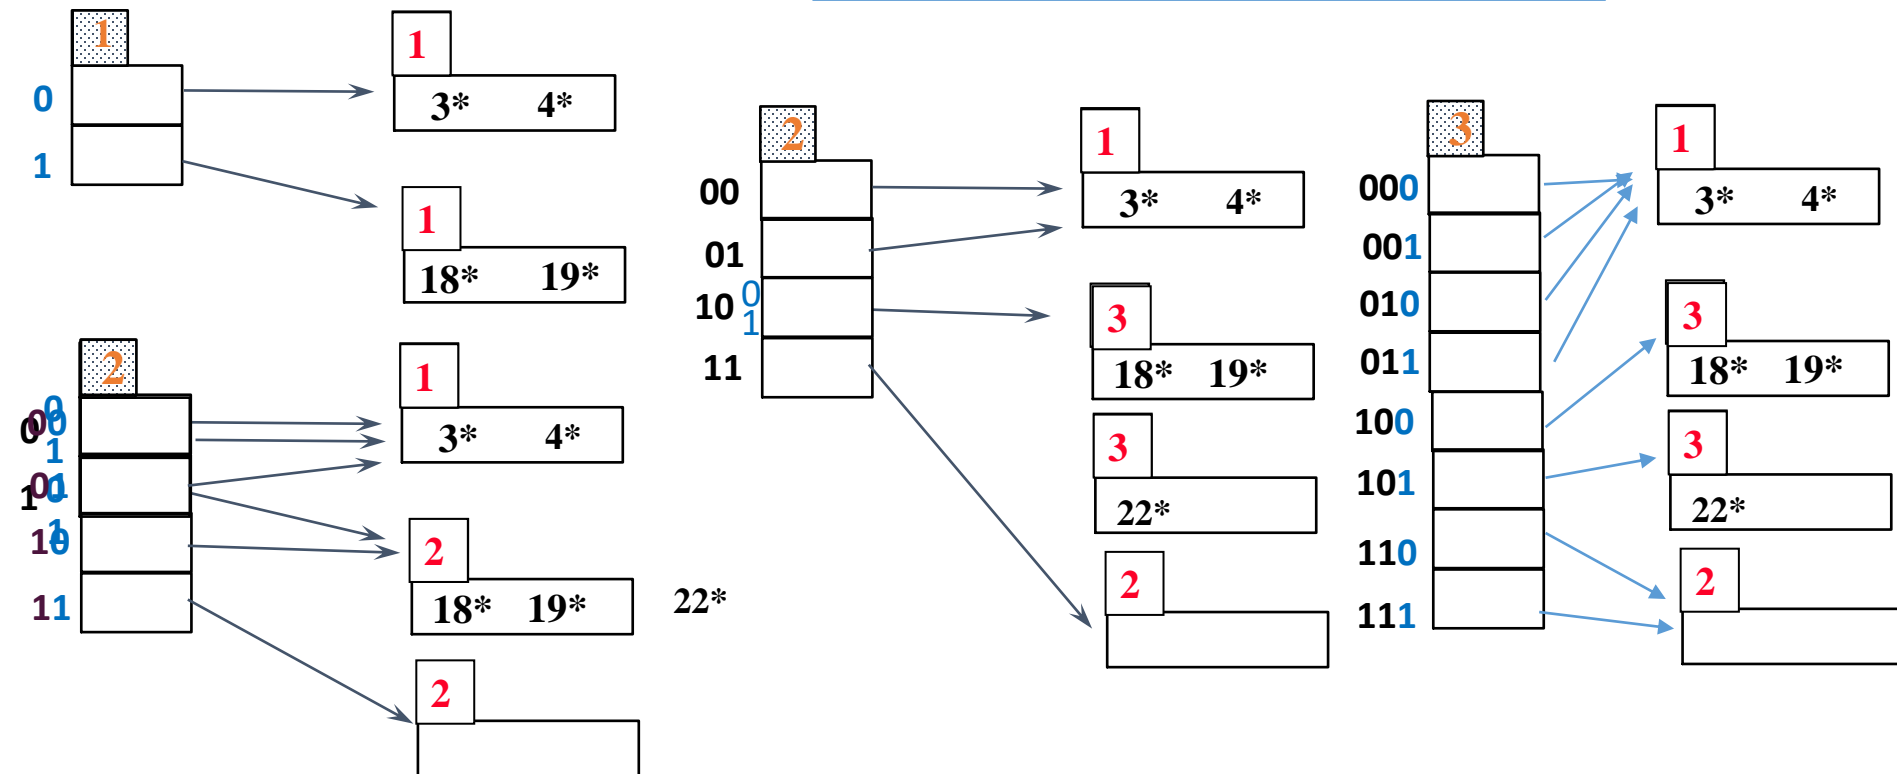
increment **local depth**

allocate new page with new **local depth**

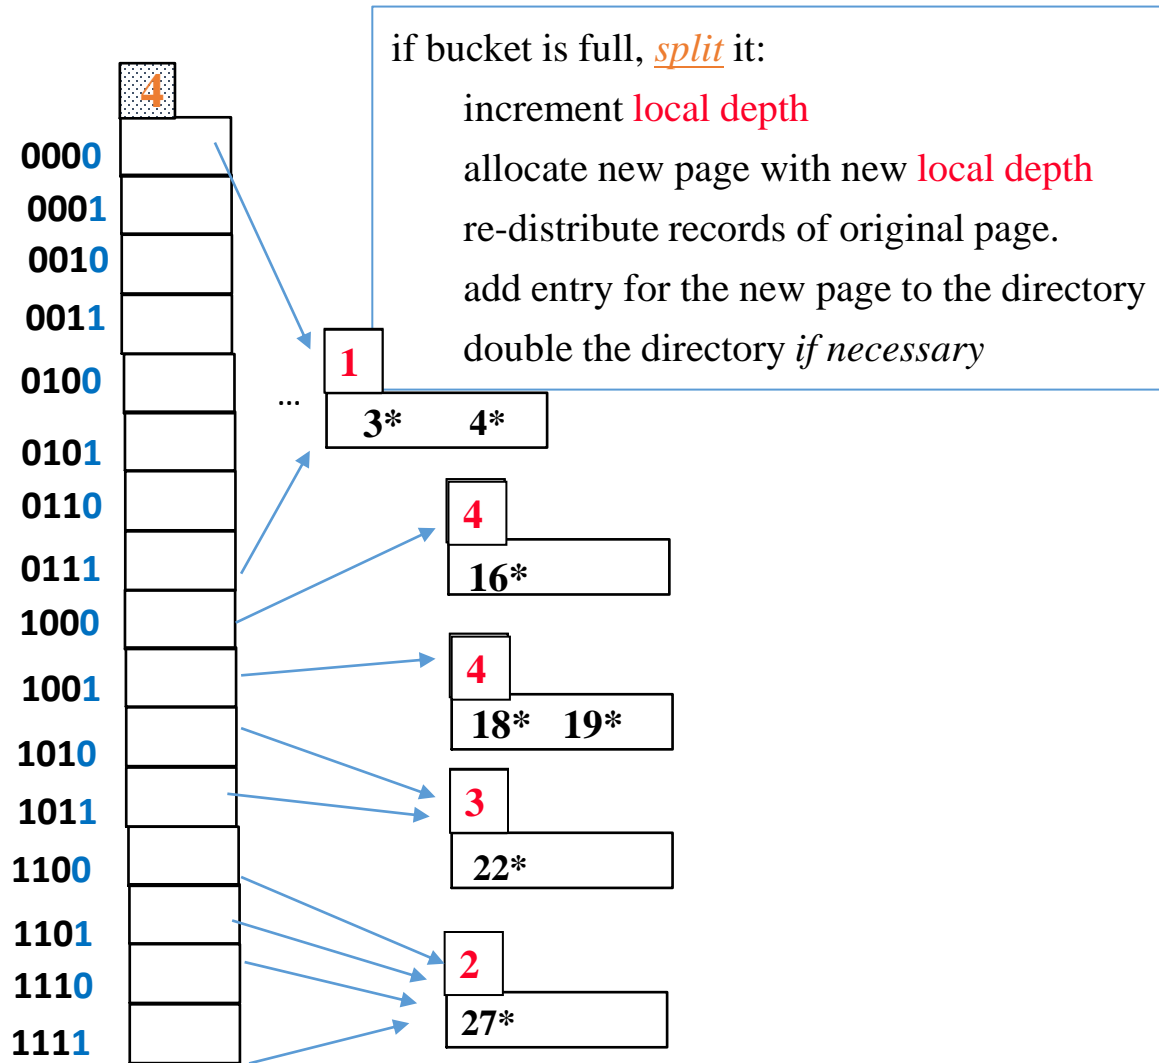
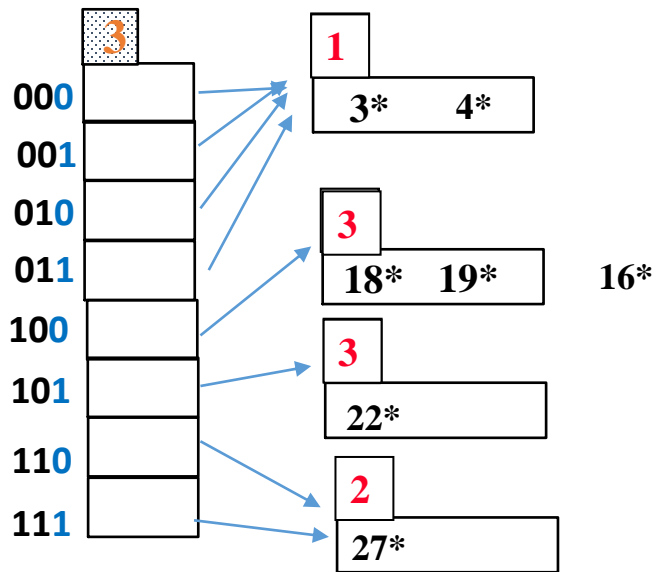
re-distribute records of original page.

add entry for the new page to the directory

double the directory *if necessary*



- (A) Key hashes to 00011 (= 3)
- (B) Key hashes to 10010 (= 18)
- (C) Key hashes to 10011 (= 19)
- (D) Key hashes to 00100 (= 4)
- (E) Key hashes to 10110 (= 22)
- (F) Key hashes to 11011 (= 27)
- (G) Key hashes to 10000 (= 16)



How many directory elements pointing to a bucket?

$2^{\text{GD-LD}}$

Summary

- Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it.
 - Directory to keep track of buckets, doubles periodically.
 - Can get large with skewed data; additional I/O if this does not fit in main memory.

Linear Hashing

Linear Hashing

- It maintains a **constant load factor**.
 - Thus, avoids reorganization.
- It does so, by **incrementally adding new buckets to the primary area**.
- In linear hashing, the **last bits in the hash number** are used for placing the data entries.

Example

e.g.

34: 100010

28: 011100

08: 001000

Desired Lf = 67% = 2/3

Last 3 bits



Lf = 14/24

= 58%

Lf = 15/24

= 63%

Lf = 16/24

= 67%

Lf = 17/24

= 70%

000

8*

16*

32*

001

17*

25*

010

34*

50*

011

11*

27*

100

28*

12*

101

5*

13*

21*

110

14*

111

55*

15*

Insert: 13, 21, 37, 12

13: 001101

21: 010101

37: 100101

12: 001100

37*		
-----	--	--

Lf = 18/24


= 0.75

Primary area

Insertion of records

- To **expand** the file: **split** an existing bucket denoted by **k digits** into two buckets using the last **k+1 digits**.

- e.g.



When to Split?

- When there are $Lf * Bkfr$ records **more than** needed for **the given Lf**.

Split when there are $Lf * Bkfr$ records **more than** needed for **the desired Lf**

Last 3 bits

Desired Lf = 67% = 2/3

$$Lf = 14 / (3 * 8)$$

$$= 58\%$$

$$Lf = 15 / 24$$

$$= 63\%$$

$$Lf = 16 / 24$$

$$= 67\%$$

000

001

010

011

100

101

110

111

8*	16*	32*
17*	25*	
34*	50*	
11*	27*	
28*	12*	
5*	13*	21*
14*		
55*	15*	

Insert: 13, 21, 37, 12

13: 001101

21: 010101

37: 100101

12: 001100

$$2/3 * 3 = \mathbf{2 \text{ records}}$$

37*		
-----	--	--

WHY?

So, we maintain a **constant load factor.**

Expanding the file

8: 001000

16: 010000

32: 1000000

$L_f = 8/27$
67%

37*		
-----	--	--

"split image" of
the bucket 000


Boundary
value



- Points to the next bucket to split
- Points to the first bucket that uses only k-digits of the hash value

0 000	16*	32*	
001	17*	25*	
010	34*	50*	
011	11*	27*	
100	28*	12*	
101	5*	13*	21*
110	14*		
111	55*	15*	
1 000	8*		

Boundary
value




0000	16*	32*	
0001	17*		
0010	34*	50*	
011	11*	27*	
100	28*	12*	
101	5*	13*	21*
110	14*		
111	55*	15*	
1000	8*		
1001	25*		
1010	26*		

k = 3

Hash # 1000: uses last 4 digits

Hash # 1101: uses last 3 digits



37*		
-----	--	--

Fetching a record

- Calculate the hash function.
- Look at the last k digits.
 - If it's less than the boundary value, the location is in the bucket labeled with the last $k+1$ digits.
 - Otherwise it is in the bucket labeled with the last k digits.
- Follow overflow chains as with static hashing.

$\text{key} < \text{bv} \rightarrow k+1 \text{ digits}$


$\text{key} \geq \text{bv} \rightarrow k \text{ digits}$


last k digits

Insertion

- Search for the correct bucket into which to place the new record.
- If the bucket is full, allocate a new overflow bucket.
- If there are now $Lf * Bkfr$ records more than needed for the given Lf ,
 - Add one more bucket to the primary area.
 - Distribute the records from the **bucket chain** at the boundary value between the original area and the new primary area buckets
 - Add 1 to the boundary value.

Bkfr = 3, desired LF= 2/3,
M=4 primary area buckets initially
What is k?



56 = 011 1000
67 = 100 0011
43 = 010 1011
79 = 100 1111
15 = 000 1111
27 = 001 1011
19 = 001 0011
64 = 100 0000

12 = 000 1100
33 = 010 0001
57 = 011 1001
65 = 100 0001
29 = 001 1101

$$LF = 2 / 3 = n / (3 * 4)$$

I reach LF at **n = 8**

$$\text{After } Lf * Bkfr = 2/3 * 3 = 2$$

more records, split!

key < bv → k+1 digits
key ≥ bv → k digits




last k digits

Bkfr = 3, desired LF= 2/3,
 4 primary area buckets initially
 k=2

56 = 011 1000
 64 = 100 0000

12 = 000 1100
 33 = 010 0001

 57 = 011 1001
 65 = 100 0001
 29 = 001 1101



After $Lf * Bkfr =$
 $2/3 * 3 = 2$ more
 records, split!

- Add one more bucket to the primary area.
- Distribute the records from the bucket chain at the boundary value
- Add 1 to the boundary value.

key < bv → k+1 digits
 key ≥ bv → k digits
 {
 last k digits

Bkfr = 3, desired LF= 2/3,
4 primary area buckets initially

$$33 = 010 \ 000\underline{1}$$

$$57 = 011 \ 1001$$

$$65 = 100 \ 0001$$

$$29 = 001 \ 1101$$

000

56*	64*	
-----	-----	--

bv → **001**

33*	57*	65*
-----	-----	-----

bv → 10

--	--	--

11

67*	43*	79*
-----	-----	-----

 →

15*	27*	19*
-----	-----	-----

100

12*		
-----	--	--

101

--	--	--

After $Lf * Bkfr = 2/3 * 3 = 2$ more records, split!

- Add one more bucket to the primary area.
- Distribute the records from the bucket chain at the boundary value
- Add 1 to the boundary value.

key < bv → k+1 digits
key ≥ bv → k digits

last k digits

Bkfr = 3, desired LF= 2/3,
4 primary area buckets initially

29 = 001 1101

2 = 000 0010

000 56* 64*

001 33* 57* 65*

bv → **0**10 2*

bv → 11 67* 43* 79* → 15* 27* 19*

100 12*

101 29*

110

After $Lf * Bkfr =$
 $2/3 * 3 = 2$ more
records, split!

Distribute 2*: stays in 010!

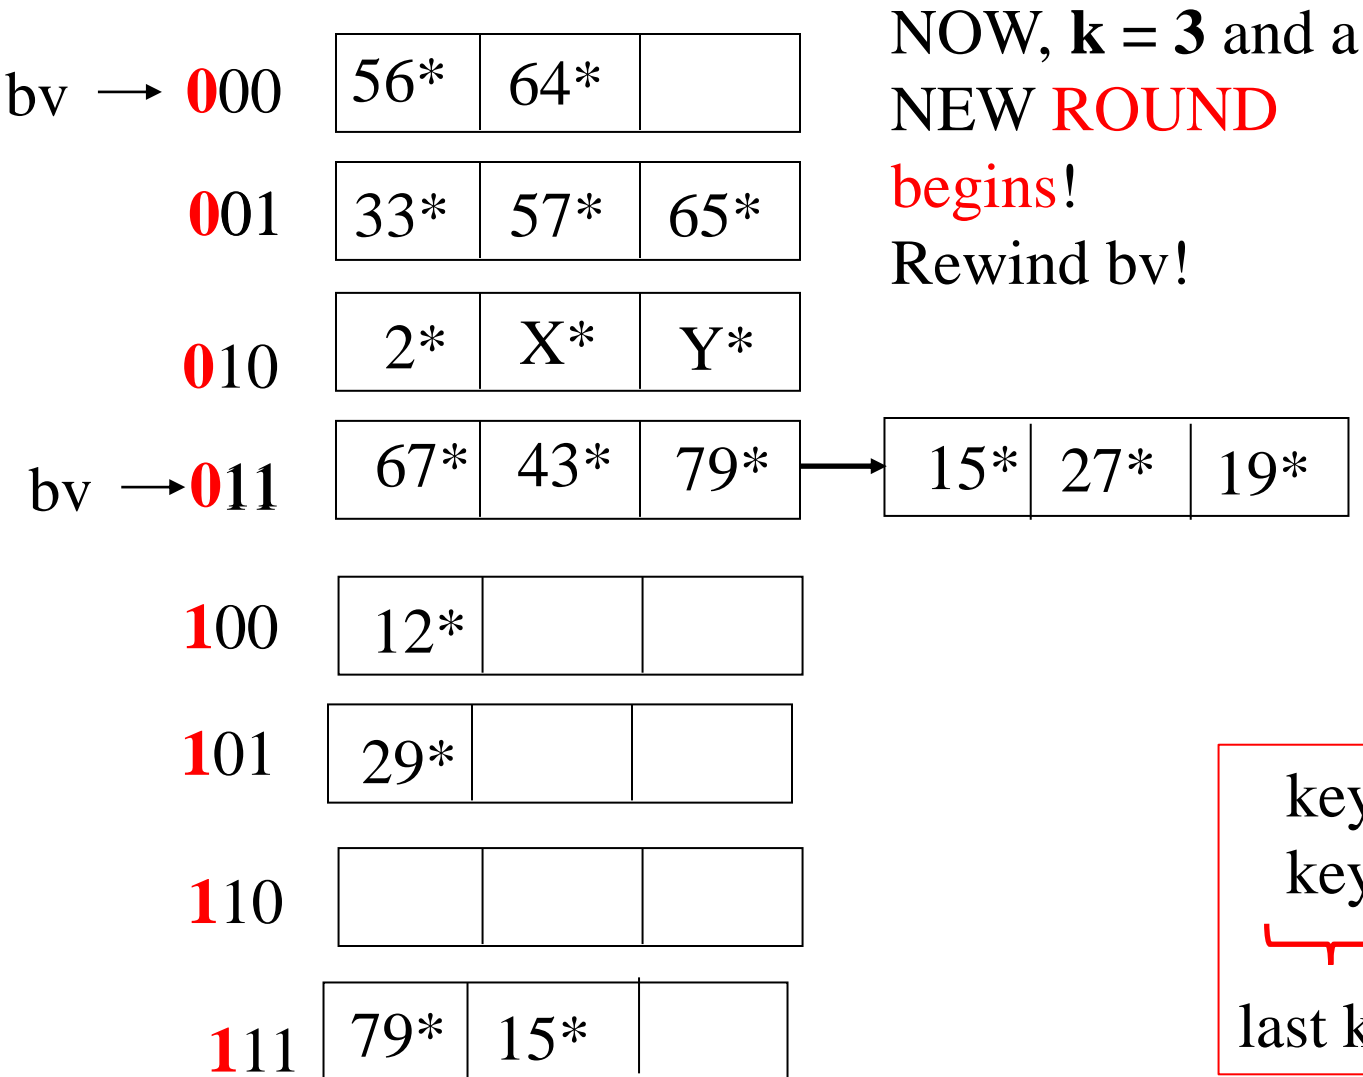
key < bv → k+1 digits

key ≥ bv → k digits

last k digits

Bkfr = 3, desired LF= 2/3,
4 primary area buckets initially

Suppose two
more recs
added



Distribute:

67 = 100 0011
 43 = 010 1011
 79 = 100 1111
 15 = 000 1111
 27 = 001 1011
 19 = 001 0011

key < bv → k+1 digits
 key ≥ bv → k digits

last k digits

bv \rightarrow **0**000

56*	64*	
-----	-----	--

bv \rightarrow 001

33*	57*	65*
-----	-----	-----

010

2*	X*	Y*
----	----	----

011

67*	43*	
-----	-----	--

\rightarrow

	27*	19*
--	-----	-----

100

12*		
-----	--	--

101

29*		
-----	--	--

110

--	--	--

111

79*	15*	
-----	-----	--

1000

--	--	--

Suppose two
more recs
added

NOW, **k = 3** and a
NEW ROUND
begins!

key < bv \rightarrow k+1 digits
key \geq bv \rightarrow k digits

last k digits, k=3 ⁶¹

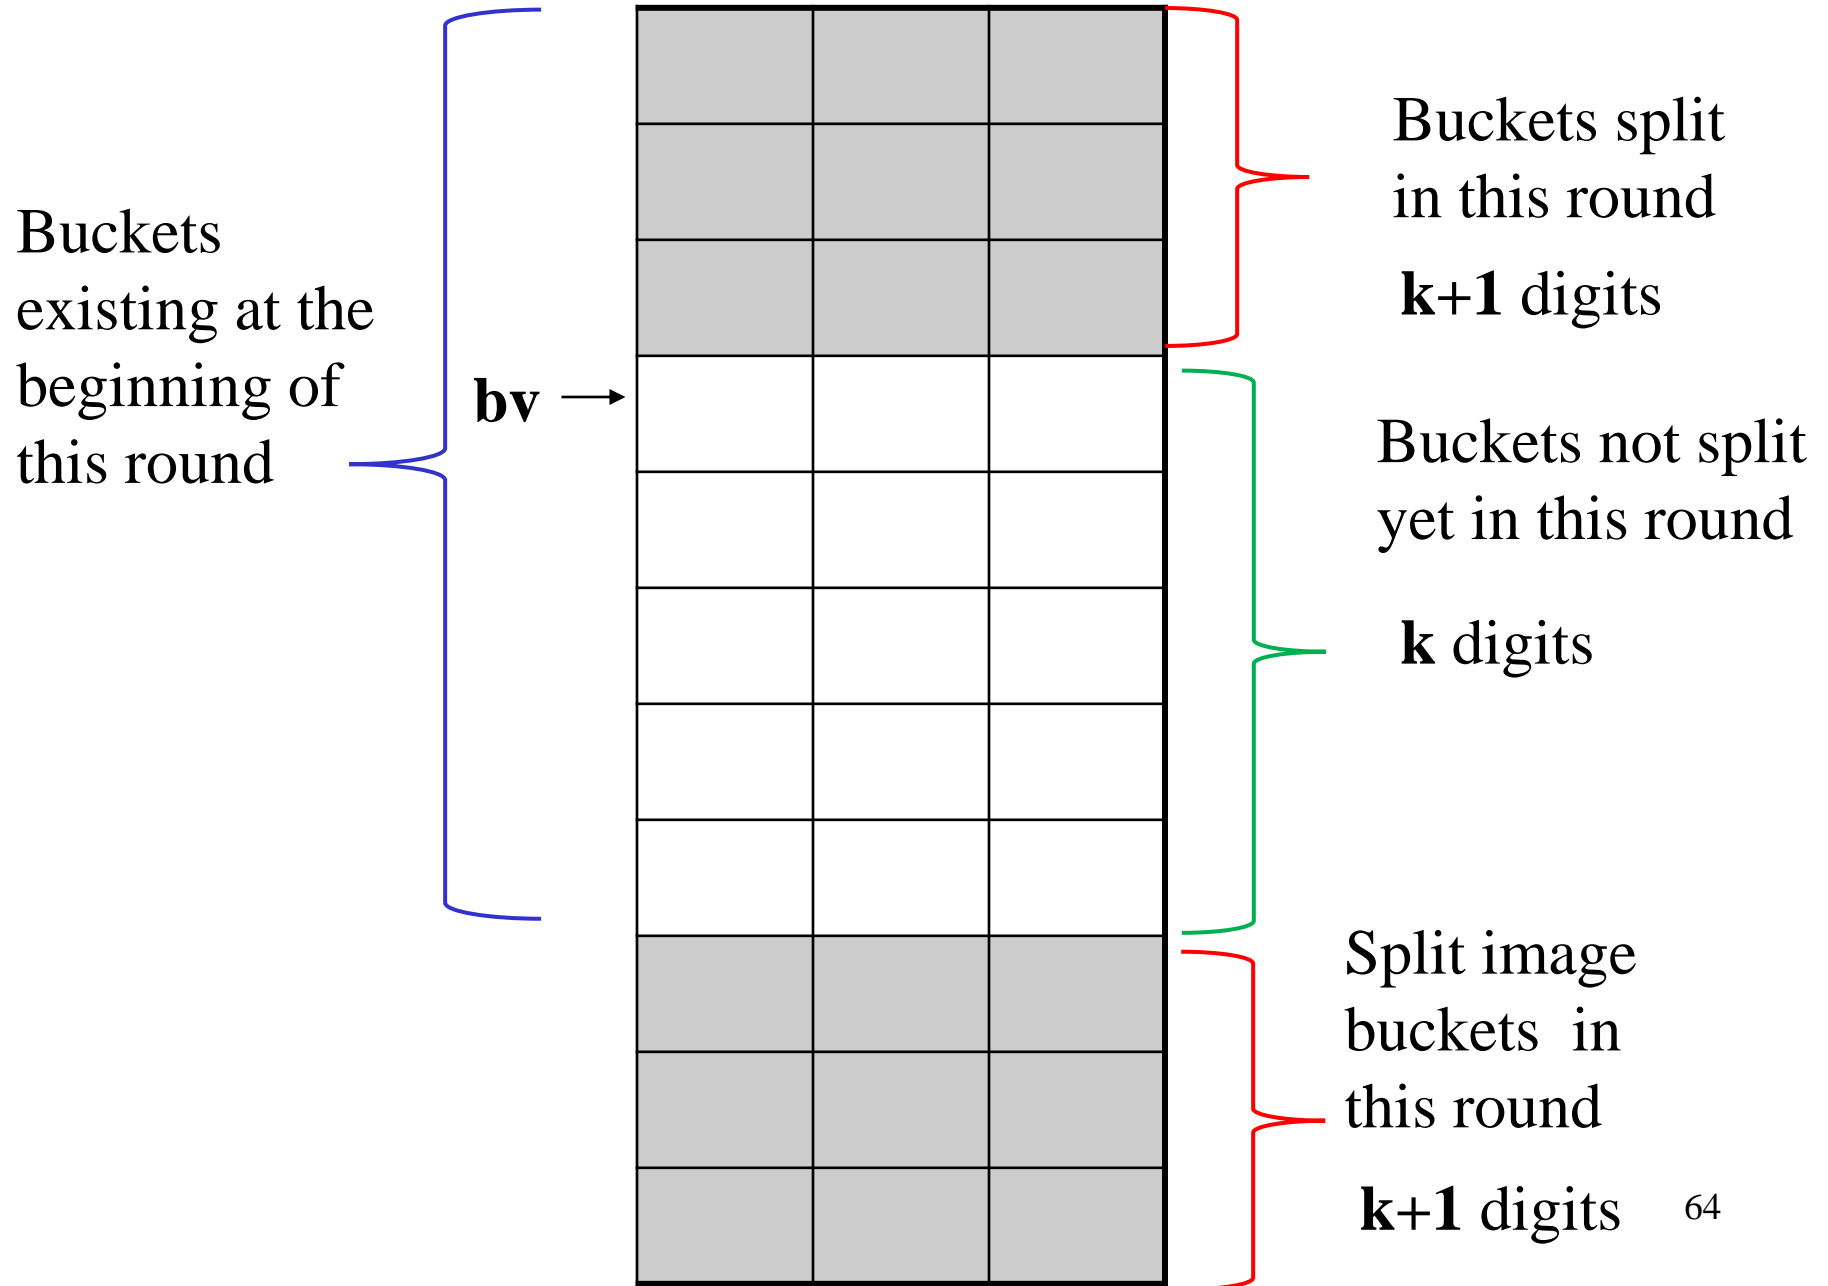
Linear Hashing

- Overflow and expansion are independent events
 - Overflow does not trigger bucket split
 - Bucket split does not necessarily remove an overflow chain (but eventually the overflowing bucket will also be split!)
- After splitting all buckets in a **round**, **rewind** **by to 0**, **increase k by 1**
 - At this point, we **doubled** the **range** into which keys are hashed!

Linear Hashing


- No reorganization
- It maintains a **constant load factor**.
- There are still overflow chains (hopefully, won't be too long, as the overflowing bucket will also be split eventually)
- Record fetch time → still close to 1 disk access

Buckets during a Round in Linear Hashing



Buckets during a Round in Linear Hashing

- Suppose we have M primary area buckets during a round that has started with k digits
 - What can you say about the quantity of M ?
 - $2^k \leq M < 2^{k+1} \rightarrow k \leq \log M < k+1$



Deletion

- Read in a chain of records.
- Replace the deleted record with the last record in the chain.
 - If the last overflow bucket becomes empty, deallocate it.
- When the number of records is $Lf * Bkfr$ less than the number needed for Lf , contract the primary area by one bucket.

Compressing the file is exact opposite of expanding it:

- Keep the total # of records in the file and buckets in primary area.
- When we have $Lf * Bkfr$ fewer records than needed, consolidate the last bucket with the bucket which shares the same last k digits.